

TFHE Notes

- Aman Savaria

Chapter 0

About the notes: Roadmap

This document contains my notes that I took when studying about the TFHE scheme on a lower level. It is divided into multiple modules/parts etc which I didnt do a good job of differentiating with chapter names. This 'disclaimer' sort of a chapter does that.

The first section spans from the first to the fifth chapter. It has an in-depth explanation of FHE encryption, arithmetic and operations such as bootstrapping. These chapters are written by following Daniel Lowengrub's blog. These chapters start off by describing the LWE ciphertext to work with scalars, their arithmetic, and then moves on to RLWE ciphertexts to work with polynomials. After that, it moves to GSW ciphertexts and talks about ciphertext-ciphertext multiplication and then describes in detail operations used in bootstrapping. It finally ends with implementing the NAND operation using bootstrapping. Note that everything that was mentioned in these chapters have been implemented in python from scratch just using numpy for better understanding. You can find the code in this repository. A small description of important terms:

- **LWE:** An LWE ciphertext encrypts an integer. It is a collection of two elements (a, b) , with a being a vector of n elements, sampled randomly from a discrete uniform distribution, and b being an integer. $b = a \cdot s + e$, where s is a vector and is called a secret key, while e is an integer called the noise and is sampled from a discrete normal distribution. This is discussed in detail more in chapter 1.
- **RLWE:** RLWE encrypts a polynomial. Like LWE, it is also a collection of (a, b) , but this time, they both are polynomials. The relationship between a and b is the same, but the noise and secret key are now polynomials instead of being an integer or a vector. This is discussed in detail in chapter 2.
- **Homomorphic NAND:** The homomorphic NAND operation is the NAND operation applied on two LWE ciphertexts, which encrypt boolean True or False. Since LWE ciphertexts encrypt integers, the boolean values are represented as 2 (True) and 0 (False). The NAND operation is evaluated by first passing the two inputs through a function that performs a linear transform to get a single LWE output, which is then fed into a bootstrapping operation to evaluate the NAND operation. This is discussed in detail in chapter 5.

The second section is just chapter 6. It was written based on the TFHE-rs handbook. This chapter goes specifically into TFHE, and describes the different kinds of ciphertexts and operations in TFHE. It has a fair bit of overlap (actually a lot of overlap) with the first section, as almost all the operations are the same and implemented more or less, similarly, except that TFHE uses the more 'general' GLWE, GLew and GGSW ciphertexts which while more complex, follow the same rules and steps for almost everything.

The third section discusses PBS (Programmable Bootstrapping) further by implementing it from scratch for the sign operation. The sign operation is defined as follows:

$$\text{sign}(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{otherwise} \end{cases} \quad (1)$$

This chapter provides a hands on example with implementation to explain how to create the TLU polynomial and then use it to evaluate a function on an LWE ciphertext using bootstrapping. It also talks about how the domain

changes if the function to be evaluated is not negacyclic and why. Finally, it talks a little about bivariate PBS, that utilizes a packing trick to pack two ciphertexts into 1 so that PBS can be implemented to evaluate functions that take two inputs while traditional PBS only takes one LWE ciphertext. This is also done with an example of showing how to use this packing trick to implement PBS to evaluate the max operation between two ciphertexts.

IMPORTANT THING TO NOTE: There might be places where I denoted a ring as $R_q(x)/(x^N + 1)$. This is a typo, they're supposed to be $\mathbb{Z}_q(x)/(x^N + 1)$ as the coefficients are integers.

0.1 Important Abbreviations/Notations

- LWE: Learning With Errors Ciphertext/Problem
- RLWE: Ring Learning With Errors Ciphertext/Problem
- PBS: Programmable Bootstrapping
- $\mathbb{Z}_q/(X^N + 1)$: A polynomial ring, where the polynomials have a max degree of N-1 and the coefficients are modulus q, $q \in \mathbb{Z}_+$.
- Cleartext: Not an abbreviation but important. A regular number before being processed/encrypted is called cleartext.
- Plaintext: Not an abbreviation but important. A plaintext is the encoded cleartext.
- TLU/LUT: A Lookup Table. Used for bootstrapping.

Chapter 1

LWE: Working With Scalars.

1.1 Introduction

1.1.1 What is special about Core-crypto?

Core crypto is the deeper low-level API for TFHE. It allows the user to control everything about encryption. In my experience, using this API is almost mandatory if you want to implement papers that work fundamentally with operations such as PBS etc. as the high level apis like FheInt16 etc. do not have everything you need to implement those as they are abstracted high level apis.

However, the api is hard to learn and very intimidating for a newbie. So Im making notes as I go for my own reference.

1.2 Encryption And Decryption

This section describes encryption and decryption for LWE ciphertexts in core-crypto, and in general. Remember that LWE is a special case of GLWE, where, if in GLWE, you have m polynomials for sk and A respectively, LWE uses m=1. This means that for LWE, b is a scalar rather than a m dimensional vector.

1.2.1 Encryption

q is called the ciphertext modulus and is a positive power of 2. This q is the *space* where you move the message and encrypt. Say you have a message $m \in \mathbb{Z}_q$, and secret key $S \in \mathbb{Z}_q^n$ (remember that this is for LWE and not for GLWE). One way to create the secret key is to have a random binary vector of size N. You create a vector **A** where each element is uniformly sampled from $\mathcal{R}_q$ (theyre all integers), and a scalar value (again, only for LWE, for GLWE, this becomes a m-dimensional vector too) e which is sampled from $\mathcal{N}(0, \sigma)$ which is called the noise. ($\sigma \ll 1$)

Now the way this noise e is actually created is:

- You first sample x from $\mathcal{N}(0, \sigma)$.
- You scale this sampled value. You want to bring it from where it is (very close to 0) to somewhere thats comparable to q (because that is where the ciphertext 'lives'). Basically, you want e to be in the range $[-\frac{q}{2}, \frac{q}{2}]$. Note that it shouldnt go beyond this range because then the magintude of the noise becomes too large to get an accurate/correct decryption.
- Now to do the above, you basically do this: $e = \lfloor \frac{q}{2} \cdot x \rfloor$

Code to create noise:

```
def generate_normal_sample(config: LWConfig):  
    '''  
    This function generates the normal sample required for encryption within the given range. Since  
    ↪ this is LWE, we only need 1 element here as m = 1.  
    '''
```

```
ei = np.int32(np.random.normal(loc=0.0, scale=config.sigma) * (config.q//2))
return ei
```

Code to create the vector a.

```
def generate_uniform_sample(config: LWEEConfig):
    """
    Generate an array of {size} elements using a uniform random distribution in the required range
    """
    # get the limits
    low = (config.q//2) * -1
    high = (config.q//2) # not subtracting -1 because high is exclusive in random.randint
    # get the random array
    a = np.random.randint(low=low, high=high, size=config.n, dtype=np.int32)
    return a
```

Code for LWE Encryption key

```
class LWEEncryptionKey:
    """
    This class implements the LWEEncryption key. The encryption key implements creating the key,
    ↪ which is just a random binary vector or n elements and encrypting a ciphertext form a
    ↪ plaintext.
    """

    def __init__(self, config:LWEEConfig):
        self.config = config
        self.key = np.ndarray # this just initializes key to be an array, which will be then
        ↪ populated

    def generate_key(self):
        """
        This function generates the secret key.
        """
        self.key = np.random.randint(low=0, high=2, size=self.config.n)
        ...
```

Now the core idea behind encryption is that you basically add noise so large to the original message that it basically overshadows the actual message. For this to happen, you want your q to be large. To achieve this, you need to limit the range of messages. You basically define an integer p , called the ciphertext modulus, which is a positive power of 2, and defines the range which limits the magnitude of the input message, meaning, $m \in [-\frac{p}{2}, \frac{p}{2}]$ (also represented as $(m \in \mathbb{Z}_q)$). Now $p < q$, and ideally, it should be considerably smaller.

Code for LWEEConfiguration:

```
class LWEEConfig:
    """
    This class defines the configurations for encryption
    """
    def __init__(self, p, q, noise_level, dimension):
        """
        p: plaintext modulus
        q: ciphertext modulus
        noise_level: The noise level (used as standard deviation for generating the noise)
        """
```

```

    dimension: m (number of elements in the sk)
    '''
    self.p = p
    self.q = q
    self.sigma = noise_level
    self.n = dimension

```

Before encryption, you encode your input message to the higher 'q' space (ill define in the next section what how is this done). And the actual message basically takes the first p MSB (most significant bits) of this q space and the rest of the space is taken by the noise. You want to have a gap between these two so that decryption can be done accurately.

Encoding

Encoding, like the name suggests is *encoding* the message $\mathbb{Z}_p$ to $\mathbb{Z}_q$. To do this, you first define, $\Delta = \frac{q}{p}$. Then you multiply: $\Delta \cdot m$. This moves the message m from $\mathbb{Z}_p$ to $\mathbb{Z}_q$. After this, we can encrypt. This encoded cleartext is also called as plaintext.¹ **Encoding code:**

```

def encode_plaintext(m, config):
    '''
    This function encodes the message before its encrypted.
    '''
    # get the delta
    delta = config.q//config.p
    encoded = np.int32(delta * m)
    return LWEPlaintext(encoded)

```

The LWEPlaintext is just a class that stores (an encoded) plaintext. Its not exactly needed but for consistency its a good idea to have it:

```

class LWEPlaintext:
    '''
    This class contains the plaintext to be encrypted.
    '''
    def __init__(self, m):
        '''
        This class (for now), just stores the message wrapping it in a class. It might contain
        ↪ code for encoding it too later but i dont know
        '''
        self.m = m

    def __str__(self):
        return str(self.m)

    def __repr__(self):
        return str(self.m)

```

Actual Encryption

First thing, you are given a message and a secret key. A good example of a secret key is an n dimensional binary vector (for LWE). You sample and create **A** and **e**. Then you define **b** such that:

$$b = \langle a, s \rangle + m \cdot \Delta + e$$

$\langle a, s \rangle$ is the dot product between a and s

¹I call it that for clarity between non-encoded and encoded.

The tuple $(\mathbf{a}, \mathbf{b}) \in (\mathcal{R}^n, \mathcal{R})$ is the LWE encryption of the message m .

Encryption Code:

```
# LWEEncryptionKey Class
def encrypt(self, plaintext:LWEPlaintext):
    '''
    Encrypt the plaintext message
    '''
    # generate a and noise
    a = generate_uniform_sample(self.config)
    e = generate_normal_sample(self.config)

    # get b
    b = np.dot(a, self.key) + plaintext.m + e

    return LWE ciphertext(a, b)
...
```

Code for LWE ciphertext class. This stores LWE ciphertext and defines arithmetic over it:

```
class LWE ciphertext:
    '''
    This class defines the ciphertext.
    '''
    def __init__(self, a, b):
        self.a = a
        self.b = b
```

Note for clarity: LWE ciphertexts encrypt a single scalar value. Therefore, in the encryption, m & e are scalars. Whereas a and s are n -dimensional vectors, where n is already specified in the encryption parameters. Remember this as this becomes a point of confusion between LWE, RLWE and GLWE ciphertexts.

1.2.2 Some information

There are actually 2 terms that are really important here, and they will be used in many places. They have a subtle difference and thus, its easy to interchange, or think that theyre the same thing, but theyre not. These are **Plaintext** and **Cleartext**. **Cleartext** is basically just your raw input, nothing done to it. **Plaintext** on the other hand, is actually the cleartext thats encoded in Z_q (by multiplying with $/\Delta$). As well see later, when working with polynomials instead of scalars, the corresponding cleartext is a polynomial $p(x)$ and its corresponding plaintext becomes a polynomial in the ring $Z_q[x]/(x^N + 1)$. This difference is very important as these terms will be used a lot and its important to be able to understand the operations well.

1.2.3 Decryption

The number (or m -dimensional vector in the case of GLWE) b contains the actual encrypted message. Decryption is basically the process of removing all the solar-system sized noise attached to it to hide m . First things first, the easiest thing to do is to remove $\mathbf{A}$ from it, as it is present as a part of the ciphertext and for decryption, the secret key is also present. So, we do:

$$b - \langle a, s \rangle = m \cdot \Delta + e$$

Decoding

From the above step, we have $m \cdot \Delta + e$, but this isnt the actual decryption, the output still contains the noise and is in $\mathcal{R}_q$, and we want the decrypted output to be in $\mathcal{R}_p$. This is where the decoding step helps (Im putting this in

decoding, but it is also a part of the decryption step, and most resources treat it as such, Im just putting it in decoding as it personally feels more like a decoding step to me as it doesnt use the secret key, but Im stupid too). First, you basically just scale this number down from $\mathcal{R}_q \rightarrow \mathcal{R}_p$. This is done by calculating:

$$d = \lfloor \frac{m \cdot \Delta + e}{\Delta} \rfloor$$

Decryption and decoding code:

```
def decode_plaintext(m, config):
    """
    This function decodes the decrypted plaintext to remove the noise and get the final output
    """
    # get the delta
    delta = config.q/config.p
    # round to remove the noise
    decoded = int(np rint(m/delta))
    # need to take a mod wrt p centered around 0 to get the final output
    final = ((decoded + (config.p/2)) % config.p) - (config.p/2)
    return final

# LWEEncryptionKey Class
def decrypt(self, cipher:LWECiphertext):
    """
    This function decrypts the ciphertext generated using the same secret key.
    """
    b = cipher.b
    a = cipher.a

    inter = b - np.dot(a, self.key) # gives m + e, and not just m
    # return LWEPlaintext(inter)

    # round and get the final decrypted output
    final = decode_plaintext(inter, self.config)

    return LWEPlaintext(int(final))
```

Now, if there is a sufficient gap between the MSBs and the noise (which is why it was important), this descaling and rounding operation brings the output back down to p (partially) and also removes the noise. This is why selecting the parameters p , q and σ correctly is important. Because otherwise, the noise might become too large and end up *mixing* with the MSBs that contain the actual message, leading to an incorrect and/or non-exact decryption. You want to make sure that:

- There is sufficient difference between the magintudes of p and q (q should be orders of magnitude larger than p).
- The value of σ should be small enough so that the noise $|e|$ does not go beyond $\frac{q}{2}$, as otherwise, it can *push* inside the MSBs.

Finally, to get the actual output, and complete the decoding process, you need to take a modulus of d wrt p , to ensure that the decryption is correct and is within the range $[-\frac{p}{2}, \frac{p}{2}]$. This can be done using this math expression (this is my favourite way, there are many like this, but this is my favourite one (I totally just stole it from a blog)):

$$d = ((d + \frac{p}{2}) \% p) - \frac{p}{2}$$

A visual explanation of why do encryption and decryption work.




Figure 1.1: Bitwise representation of LWE/RLWE/GLWE ciphertext.

Any FHE ciphertext (LWE, RLWE, GLWE etc. and other types) can be represented as shown in the figure above. The ciphertext modulus (q) defines this number of bits (each bit is a small square in the figure). The blue squares are where the message is stored and as it can be seen, it's stored in the MSBs (Most Significant Bits) here and its length is defined based on the value of plaintext modulus (p), while the noise (the red squares) is stored in the LSB (Least significant bits). Having $p = 8(2^3)$ and $q = 2^{32}$ means that that message would be stored in the first 3 bits out of the 32 bits which would be the length of the entire ciphertext.

This is done by design. As defined above, during encryption of the plaintext by doing $m \cdot \Delta$, where $\Delta = \frac{q}{p}$, we are basically 'moving' our message from the $\log_2 p$ bit message to $\log_2 q$ message length, with the actual message taking the $\log_2 p$ MSBs. All the other values like q , e etc are also in the q space as well. The way the noise is constructed is such that it will always take the LSB spaces, as we take samples for a normal distribution such that these samples are $\ll 1$, which are then multiplied with q , which, while elevating the noise to this q space, is still much smaller than the message encoding. Therefore, when they are added, we naturally get this structure shown in the figure above.

The separation seen between the red and blue bits is very important for decryption to be accurate and exact. During decryption and decoding, the result of $b - \langle a, s \rangle$ is rounded down to the p space. If the noise and actual message bits are not 'separated' like this, the rounding will not be accurate and will contain traces of noise, making decryption inaccurate. This is also the reason why you can't do a lot of chained addition and multiplication (multiplication even moreso) with ciphertexts as they increase the noise and it starts moving left.

Following needs to be cared for to ensure this 'gap':

- Ensure that $q \gg p$.
- The standard deviation σ should be small enough that the noise does not become too large (remember that we are multiplying the normal sample with q), in my experiments, where $p = 8$ and $q = 2^{32}$, I have used $\sigma = 2^{-20}$. However, σ should not be so small that the noise becomes extremely small as that compromises with the security. (Note that I'm not sure if the σ that I've used is secure or no, I used it from a tutorial).

1.2.4 Notes about encryption and decryption in Core-Crypto

The core-crypto API in rust requires you to do a few shenanigans (technically they're the 'correct' ways to do something like but I'm a python pleb so I'll call it a shenanigan). The API actually works with only unsigned integers and thus, you can't just work with a negative number as is. You need to cast any signed integer (positive or negative) as the corresponding unsigned integer. Also, during decryption, in the decoding process, after performing mod wrt the plaintext modulus, do not forget to cast it back to the signed integer to get the correct answer.

1.3 Arithmetic

This section talks about how arithmetic is done in $LWECiphertexts$. The following operations are possible:

- Addition: Ciphertext-ciphertext and Ciphertext-Plaintext addition.
- Subtraction: Same as addition.
- Multiplication: Ciphertext-plaintext multiplication. Multiplication between two ciphertexts isn't possible it seems.

1.3.1 Addition

An LWE ciphertext is defined as (a, b) , where $a \in \mathcal{Z}^n$ and $b \in \mathcal{Z}$. Since both a and b are going to be random, we denote by $LWE_s(m)$, the family of ciphertexts that encrypt the cleartext message m under the secret key s .

Ciphertext-Plaintext addition

Say you want to add a ciphertext encrypting a cleartext message m , $L_1 = LWE_s(m) = (a, b) \in (\mathcal{Z}^n, \mathcal{Z})$, and a scalar, constant value c . We denote this operation by $PAdd$, such that $PAdd(L_1, c) = LWE_s(m + c)$. To add c to L_1 , you do the following:

1. Encode c by doing $\Delta \cdot c$.
2. Add $\Delta \cdot c$ to b of L_1 .
3. The resulting ciphertext $(a, b + c)$ is the answer.

Why does this work?

To decrypt, $L_1 = (a, b)$ you do:

$$b - \langle a, s \rangle = \Delta \cdot m + e.$$

Followed by $\lfloor \frac{\Delta \cdot m + e}{\Delta} \rfloor$ which gives the decrypted result m back.

Similarly, for $LWE_s(a, b + \Delta \cdot c)$:

$$(b + \Delta \cdot c) - \langle a, s \rangle = \Delta \cdot m + \Delta \cdot c + e$$

Followed by $\lfloor \frac{\Delta(c+m)+e}{\Delta} \rfloor$ which will give the decrypted result $(m + c)$ back.

Ciphertext-Ciphertext Addition

Consider two ciphertexts $L_0 = LWE_s(m_0)$ and $L_1 = LWE_s(m_1)$ encrypted under the same encryption key s . You want to define a function $CAdd(L_0, L_1) = L_{sum}$ such that $Dec(L_{sum}) = (m_0 + m_1)$. If the ciphertexts are $L_0 = (a_0, b_0)$ and $L_1 = (a_1, b_1)$, then $L_{sum} = (a_0 + a_1, b_0 + b_1)$.

Why does this work?

Decrypting L_{sum} :

$$(b_0 + b_1) - (a_0 + a_1)s = \Delta(m_0 + m_1) + (e_0 + e_1)$$

Dividing this by Δ and rounding will give the decrypted $m_0 + m_1$.

Note on the noise increase: As you can observe, the noise in the resulting ciphertext is higher as it is the sum of noise from both the added ciphertexts. This is expected, and while it does not increase noise a lot (as it's only adding), you can still not do a lot of chained additions because after a point the noise becomes too large for proper decryption to happen. However, if your parameters are set correctly, this point is so far away that it is not a point of worry for most applications.

1.3.2 Subtraction

Subtraction works the same way you just replace the $+$ by the $-$.

Multiplication

For LWE ciphertexts, only multiplication with a scalar is supported. Consider a ciphertext $L = LWE_s(m) = (a, b)$ encrypting a message m , which is to be multiplied with a cleartext scalar c . We define a function $PMul(L, c) = LWE_s(c \cdot m)$, where $PMul(L, c) = (a \cdot c, b \cdot c)$.

Why does this work?

Decrypting $PMul(L, c)$:

$$\begin{aligned} b \cdot c - a \cdot c \cdot s &= c(m + e) \\ b \cdot c - a \cdot s \cdot c &= (c \cdot m) + (c \cdot e) \end{aligned}$$

Decoding this by dividing by Δ and rounding will give $m \cdot c$. **Code for LWE Ciphertext arithmetic:**

```
# LWE Ciphertext Class
def add_ciphertext(self, other: LWE Ciphertext):
    """
    This function adds a given ciphertext to the current ciphertext and returns a new
    ↪ ciphertext.
    """
    # adding their as and bs
    added_a = np.add(self.a, other.a, dtype=np.int32)
    added_b = np.add(self.b, other.b, dtype=np.int32)

    ## taking mods just to see if this works, but we have to take mod wrt q
    # q = config.q
    # added_a = ((added_a + (q//2)) % q) - (q//2)
    # added_b = ((added_b + (q//2)) % q) - (q//2) # im getting overflow errors when doing
    ↪ this

    return LWE Ciphertext(added_a, added_b)

def sub_ciphertext(self, other: LWE Ciphertext):
    """
    This function subtracts the other ciphertext from the current one.
    """
    sub_a = np.subtract(self.a, other.a, dtype=np.int32)
    sub_b = np.subtract(self.b, other.b, dtype=np.int32)
    return LWE Ciphertext(sub_a, sub_b)

def mul_plaintext(self, other: np.int32):
    """
    Multiply the other integer to current ciphertext
    """
    # we just need to multiply the give plaintext to both a and b
    mul_a = np.multiply(other, self.a, dtype=np.int32)
    mul_b = np.multiply(other, self.b, dtype=np.int32)
    return LWE Ciphertext(mul_a, mul_b)

def add_plaintext(self, other: LWE Plaintext):
    """
    Add the other plaintext to the current ciphertext. This function assumes that the input
    ↪ plaintext message is already encoded to R_q.
    """
    # we just need to scale other and then add it to b
    addb = np.add(self.b, other.m, dtype=np.int32)
    return LWE Ciphertext(self.a, addb)
```

1.4 Trivial LWE Ciphertext

This is just a ciphertext (a, b) where a is a vector of zeros and b is just the encoded message with no noise added to it. It will be used down the line to enable RLWE ciphertext arithmetic.

Chapter 2

Working With Polynomials.

2.1 Why do we need to use Polynomials?

So from what I've understood in my limited reading so far, whenever you do any operation in any kind of a FHE/HE scheme, including TFHE, the noise e increases, it does so especially during multiplication. Now, that makes it not very practical when we want to chain operations together, and do operations that are more complex than regular arithmetic to implement complex (relatively speaking) functions. However, there is another operation, called *Bootstrapping*, which will be discussed in later chapters, which basically 'resets' this noise, thus allowing for chaining of complex operations in TFHE. Now for doing this bootstrapping operation, we want to work with polynomials as polynomials make it easy for us to implement bootstrapping.

2.2 Polynomial Ring and Negacyclic Rings

These are two terms that are used quite a lot and also form the basis of working with polynomials in TFHE. A polynomial ring is described/shown as follows:

$$R_q[x]/(x^N + 1)$$

There are two separate things here, the first is $R_q[x]$, and the second is $(x^N + 1)$. $R_q[x]$ means that for any polynomial that belongs to this ring, all of its coefficients are mod q . Meaning, if you define any polynomial as $p(x) = c_0 + c_1x + c_2x^2 + \dots + c_nx^n$, its coefficients $[c_0, c_1, c_2, \dots, c_n] \in R_q$, meaning they are in the range $[-\frac{q}{2}, \frac{q}{2})$. The $(x^N + 1)$ plays the same role here as q played for $R_q[x]/q$, but for the degree of polynomial this time, the polynomial that belongs in this ring is considered modulo $(x^N + 1)$. To make a polynomial modulo $(x^N + 1)$, you just divide it with $(x^N + 1)$ and use the remainder. A polynomial can have an arbitrarily high degree, but taking a mod caps it to be less than N .

Now why do we call it a ring? Like in a scalar ring, $\mathbb{Z}/q\mathbb{Z}$, where two scalars are equal modulo q if they differ by a multiple of q . Here as well, two polynomials are considered to be equal modulo $(x^N + 1)$ if they differ by a multiple of $(x^N + 1)$.

(Important Note: This distinction is something that I 'observed' from the literature that I read, it may or may not be exactly the truth (NOTE TO SELF: Dig deeper), but nonetheless, I think its better to use this nomenclature as it helps differentiate.)

Coming to the *negacyclic* thing, a polynomial who is modulo $(x^N + 1)$, basically 'loops' around to 1 once it's degree crosses N , but with a negative sign. For example, if you have $N = 4$, so $(x^N + 1)$ becomes $(x^4 + 1)$, and you have a polynomial $1 + x^5$, then:

$$\begin{aligned} 1 + x^5 &= 1 - x \pmod{(x^N + 1)} \quad (1 - x \text{ is the remainder.}) \\ x^8 + x^5 &= x^2 - x \pmod{(x^N + 1)} \end{aligned}$$

As you can see, that as the degree of the polynomial got larger than N , the powers of x that were larger than N (4) basically just wrapped around to one, but with an opposite sign. (x^5 became x , x^8 became x^2). This is what negacyclic means, and this is the behavior of polynomials in a ring. In fact, this can also be used as a "trick" to calculate a polynomial modulo $(x^N + 1)$, you just replace all instances of x^N by -1 , and you will get the modulo N

version of that polynomial. Note that for any operations done between two polynomials that are in the ring, the result will also be modulo $(x^N + 1)$. Example: Consider two polynomials $f(x) = 1 + x^3$ and $g(x) = x^2$, and they're in the ring $\mathbb{Z}_q[x]/(x^4 + 1)$, then $f(x) \cdot g(x)$ will be:

$$f(x) \cdot g(x) = x^2 + x^5 = x^2 - x \pmod{(x^N + 1)}$$

One very important thing to realize and understand is that when we say that a polynomial being in a ring $\mathbb{Z}_q[x]/(x^N + 1)$, it doesn't mean that it can not have a degree $> N$, it can. However, the process of doing modulus $x^N + 1$ the degree gets reduced to a number that smaller than N . Try to relate it to plaintext modulus, when we say that we have a number x s.t. $x \in R_q$, it also means that any number that is beyond the range $[-\frac{q}{2}, \frac{q}{2})$, can be represented as another number that is in that range. Meaning, integers that are within that range, sort of act as a 'representative integer' for everything outside of that range who have the same mod q . Similarly, any polynomial $f(x)$ that is in the ring, sort of represents a family of polynomials that have the same modulus $\mathbb{Z}_q[x]/(x^N + 1)$. This is a very subtle thing that I understood and is something that is important conceptually. So whenever we talk about a polynomial $f(x) \in \mathbb{Z}_q[x]/(x^N + 1)$, $f(x)$ really represents a family of polynomials that would share the same mod, and not just a specific polynomial.

```
class Polynomial:
    """
    This class represents a polynomial in the ring  $\mathbb{Z}_q[x]/(x^N + 1)$ 
    Note that the coefficients here are in increasing order of degrees of  $x$ , [constant,  $x$ ,  $x^2$ ,
    ↪ ...,  $x^{N-1}$ ]
    """

    def __init__(self, N, coeffs: np.ndarray):
        self.N = N
        self.coeffs = coeffs

    def __str__(self):
        """
        Print the ciphertext as an equation.
        """
        prtstr = ""
        nonzerocnt = 0
        for i, c in enumerate(self.coeffs):
            if c:
                if nonzerocnt:
                    prtstr += " + "
                prtstr += f"{c}x^{i}"
                nonzerocnt += 1

        return prtstr

    def polynomial_const_multiply(self, c: np.int32):
        """
        Multiply the current polynomial with the given constant integer and return it as a new
        ↪ polynomial.
        """
        new_p = Polynomial(self.N, np.multiply(c, self.coeffs, dtype=np.int32))
        return new_p

    def polynomial_multiply(self, p2: Polynomial):
        """
        Multiply the current polynomial with the given polynomial and return the new polynomial.
        """
        # get the N from p2
        p2N = p2.N
        # multiply and pad the results to have 2N-1 (thats the longest polynomial you can get)
```

```

    # the np.polymul function requires the coefficients in decreasing order of powers of x,
    ↪ se we reverse the current coeff array and then reverse the final output's back
    ↪ again.
    poly_prod = np.polymul(self.coeffs[::-1], p2.coeffs[::-1])[::-1]
    # creating the final answer with padded coeffs
    poly_padded = np.zeros(2*self.N - 1, dtype=np.int32)
    poly_padded[:poly_prod.shape[0]] = poly_prod

    # now, taking modulus of the result (poly_padded) wrt  $x^N + 1$ 
    result = poly_padded[:self.N]
    result[:-1] -= poly_padded[self.N:]
    return Polynomial(self.N, result)

def polynomial_add(self, p2: Polynomial):
    """
    This function adds a polynomial to the given polynomial and returns the output as a new
    ↪ polynomial.
    """
    add_coeffs = np.add(self.coeffs, p2.coeffs, dtype=np.int32)
    return Polynomial(self.N, add_coeffs)

def polynomial_subtract(self, p2: Polynomial):
    """
    This function subtracts the given polynomial from the current one and returns the new
    ↪ polynomial.
    Note: Its self - p2 and not p2 - self.
    """
    sub_coeffs = np.subtract(self.coeffs, p2.coeffs, dtype=np.int32)
    return Polynomial(self.N, sub_coeffs)

def make_zero_polynomial(N:int):
    """
    Declare a zero polynomial and return it. Basically a polynomial where all the coefficients are
    ↪ zeros.
    """
    return Polynomial(N = N, coeffs = np.zeros(N, dtype=np.int32))

def build_monomial(N, i, c):
    """
    This function builds a monomial  $c * x^i$  in the ring  $R[x]/(x^N + 1)$ 
    NOTE: Rememeber that the coeffs are in increasing order of powers of x. [constant, x, x**2, ...,
    ↪ x**N]
    """
    coeffs = np.zeros(N, dtype=np.int32)

    # find k such that  $0 \leq i + kN < N$ 
    i_mod_n = i % N
    k = (i_mod_n - i)//N

    # if k is odd then the monomial gets a negative sign because
    #  $x^i = (-1)^k + x^{(i + kN)} = (-1)^k + x^{(i \% N)}$ 
    sign = 1 if k % 2 == 0 else -1
    coeffs[i_mod_n] = sign * c

    return Polynomial(N, coeffs)

```

2.3 RLWE Ciphertexts

Like LWE Ciphertexts encrypt scalar values, RLWE Ciphertexts encrypt polynomials. Meaning, the following things change in RLWE Ciphertexts compared to LWE Ciphertexts.

- The secret key here becomes a polynomial with degree $N-1$ instead of a vector. Again, one way of creating a secret key can be having a polynomial with binary coefficients, meaning, $s(x) = s_0 + s_1x + s_2x^2 + \dots + s_{N-1}x^{N-1}$.
- a is now a polynomial too, instead of a vector.
- b is a polynomial instead of a scalar value.

However, the encryption process remains the same. Also, the configurations, p (plaintext modulus), q (ciphertext modulus) and the std deviation (σ) remain the same. The parameter N which was used to define the length of the encryption key and a , are now used as the degree of the secret key $s(x)$ and $a(x)$ polynomials.

```
class RLWEConfig:
    """
    This class defines the RLWE config.
    """
    def __init__(self, N, p, q, sigma):
        self.N = N
        self.p = p
        self.q = q
        self.sigma = sigma
```

```
class RLWEEncryptionKey:
    """
    This class defines the RLWE encryption/secret key and performs encryption and decryption of
    ↪ ciphertexts.
    """
    def __init__(self, config:RLWEConfig):
        self.config = config
        self.key = Polynomial # this time, the key is also a polynomial.

    def generate_key(self):
        """
        Generate the key. The key is a random polynomial of degree N with binary coefficients.
        """
        self.key = Polynomial(
            N=self.config.N,
            coeffs = np.random.randint(
                low = 0,
                high = 2,
                size = self.config.N,
                dtype = np.int32
            )
        )
```

2.3.1 Encoding

To encode a cleartext polynomial into plaintext, you encode the individual coefficients the same way a scalar is encoded for LWE. After encoding a polynomial $p(x) = c_0 + c_1x + c_2x^2 + \dots + c_{N-1}x^{N-1}$ becomes $p'(x) = c'_0 + c'_1x + c'_2x^2 + \dots + c'_{N-1}x^{N-1}$, where $c'_i = c_i \cdot \Delta$; $\forall i \in [0, N-1]$ where $\Delta = \frac{q}{p}$.


```
def encode_rlwe(p: Polynomial, config: RLWEConfig):
    """
    This function encodes the cleartext polynomial and returns it as a Plaintext.
    """
    # get the delta
    delta = config.q//config.p
    # scale each coefficient of the polynomial by multiplying it with delta
    scaled_coeffs = np.multiply(p.coeffs, delta, dtype=np.int32)
    return RLWEPlaintext(Polynomial(p.N, scaled_coeffs), config=config)
```

```
class RLWEPlaintext:
    """
    This class implements the RLWE Plaintext. This just stores the encoded message polynomial.
    ↪ Functionally, it isnt really needed but Ive just added it to make my code consistent with
    ↪ how TFHE libraries work.
    """
    def __init__(self, p: Polynomial, config: RLWEConfig):
        self.p = p
        self.config = config
```

2.3.2 Encryption

An RLWE encryption of a polynomial $p(x)$ under the secret key $s(x)$ is a tuple $(a(x), b(x))$. Such that:

$$b(x) = a(x) \cdot s(x) + p'(x) + e(x)$$

where $a(x) \cdot s(x)$ is the polynomial multiplication between $a(x)$ and $s(x)$ mod $(x^N + 1)$. $e(x)$ is the noise polynomial whose coefficients are created in the same way the noise e was created for LWE ciphertexts. $a(x)$ is also a polynomial whose coefficients are created the same way the vector a was created for LWE ciphertexts. Note that the $p'(x)$ here is the encoded polynomial (plaintext) and not the raw polynomial (cleartext). Now since RLWE encryptions are random, as elements of it such as $a(x)$ and $e(x)$ are generated randomly, a family of RLWE ciphertexts that encrypt a polynomial $p(x)$ under the secret key $s(x)$ is denoted by: $Enc_{s(x)}^{RLWE}(p(x)) = \langle a(x), b(x) \rangle$. Note that sometimes $RLWE_{s(x)}(p(x))$ might be used to refer to them as well just because its easier and faster to write.

```

def rlwe_generate_uniform_sample(config: RLWEConfig) -> Polynomial:
    """
    This function generates a polynomial of the degree specified in the config where each
    ↪ coefficient is an integer sampled from a uniform random distribution.
    """
    # get the limits
    minval = -config.q//2
    maxval = config.q//2

    # generate N values between the min (inclusive) and max (exclusive)
    coeffs = np.random.randint(minval, maxval, size=config.N, dtype=np.int32)

    return Polynomial(config.N, coeffs=coeffs)

def rlwe_generate_normal_sample(config: RLWEConfig) -> Polynomial:
    """
    This function generates a polynomial of the degree as per the config where each coefficient is
    ↪ drawn from a normal distribution with mean 0 and standard distribution given in the config.
    """
    coeffs = np.int32(np.multiply(np.random.normal(0, config.sigma, size=config.N), config.q//2))
    return Polynomial(config.N, coeffs)

```

```

# RLWECiphertext Class...
def encrypt(self, mx: RLWEPlaintext):
    """
    Encrypt the input Polynomial p into an RLWE Ciphertext.
    """
    # generate a and noise e
    a = rlwe_generate_uniform_sample(self.config)
    e = rlwe_generate_normal_sample(self.config)

    # encrypting
    a_times_s = a.polynomial_multiply(self.key) # <a, s>
    as_plus_m = a_times_s.polynomial_add(mx.p) # <a, s> + m (m is already encoded, which is why I
    ↪ wrote m and not delta * m)
    b = as_plus_m.polynomial_add(e)

    return RLWECiphertext(a, b, self.config)
...

```

```

class RLWECiphertext:
    """
    This class implements the RLWE ciphertext.
    """
    def __init__(self, a: Polynomial, b: Polynomial, config: RLWEConfig):
        self.config = config
        self.a = a
        self.b = b

```

2.3.3 Decryption

The decryption also is similar to the LWE decryption. You first do:

$$b(x) - a(x) \cdot s(x) = p'(x) + e(x)$$

After that, you round and decode the $p'(x) + e(x)$ to get $p(x)$.

Decoding

To decode the obtained $p'(x) + e(x)$, you just do $\frac{p'(x)+e(x)}{\Delta}$ to finally get $p(x)$.

```
def decode_coeff(c, delta, p):
    """
    This function decodes a specific coefficient within the polynomial.
    """
    # divide by delta and round
    decoded_c = int(np rint(c/delta))
    # calculate decoded_c mod p and you get the final decoded value
    out = ((decoded_c + (p//2)) % p) - (p//2)
    return out

def decode_rlwe(p: Polynomial, config: RLWEConfig):
    """
    This function decodes the given polynomial to remove the noise and get the final decrypted
    ↪ cleartext.
    """
    # calculate the delta
    delta = config.q//config.p
    raw_coeffs = p.coeffs
    # decoded each coeff and get the output polynomial
    decoded_coeffs = np.array([decode_coeff(c, delta, config.p) for c in raw_coeffs])

    return Polynomial(config.N, decoded_coeffs)
```

```
# RLWEEncryptionKey Code...
def decrypt(self, cx: RLWECiphertext) -> Polynomial:
    """
    Decrypt the given ciphertext, decode it and return the cleartext polynomial.
    """
    # separate a and b
    ca = cx.a
    cb = cx.b
    # subtracting <a, s> from b
    subbed = cb.polynomial_subtract(ca.polynomial_multiply(self.key))
    # decode: divide by delta to remove the noise and take mod wrt p
    decoded = decode_rlwe(subbed, self.config)

    return decoded
```

2.3.4 Switching Between LWE and RLWE Encryption Keys

Remember that LWE encryption key is a vector of integers (for example, binary values), while RLWE encryption key is a polynomial with integer coefficients (binary coefficients for instance). It is possible and very trivial to switch an LWE key to RLWE key. To do this, you just use the N length vector key from LWE as coefficients for $s(x)$, the RLWE key. It is important to remember that this will only work and should be done only when the configuration parameters for both the schemes are exactly the same.

```
def lwe_to_rlwe_key(lwe_key: LWEEncryptionKey, config: RLWEConfig):
    """
    This function creates an RLWE key for/from the given RLWE key.
    """
    # creating a polynomial with the coeffs as the lwe_key
```

```

if lwe_key.config.n != config.N:
    raise ValueError("The size of LWEKey and degree of RLWEConfig do not match.")
key_poly = Polynomial(N=config.N, coeffs=lwe_key.key)
rlwe_key = RLWEEncryptionKey(config=config)
rlwe_key.key = key_poly

return rlwe_key

```

2.3.5 Arithmetic

Again, just like LWE ciphertexts, there are three types arithmetic operations that can be done with RLWE ciphertexts:

- Ciphertext Addition and Subtraction: CAdd, and CSub
- Plaintext Addition and Subtraction: PAdd and PSub
- Plaintext - Ciphertext Multiplication: PMul
- Multiplying two RLWE ciphertexts is not trivially possible. There are ways to do ciphertext multiplication, between an RLWE ciphertext and something called a GSW ciphertext, which will be discussed in the coming chapters.

1. Ciphertext Addition and Subtraction

This section primarily describes how ciphertext addition is done. Subtraction is done the same way, except there you subtract instead of add (sounds silly but will make sense soon). Adding two ciphertexts $R_1 = RLWE_{s(x)}(f_1(x))$ and $R_2 = RLWE_{s(x)}(f_2(x))$ encrypting $f_1(x)$ and $f_2(x)$ is denoted by $CAdd(R_1, R_2)$ and it returns an RLWE ciphertext encrypting $f_1(x) + f_2(x)$.

$$R_1 = \langle a_1(x), b_1(x) \rangle \text{ and } R_2 = \langle a_2(x), b_2(x) \rangle$$

$$CAdd(R_1 + R_2) = \langle a_1(x) + a_2(x), b_1(x) + b_2(x) \rangle = RLWE_{s(x)}(f_1(x) + f_2(x))$$

Similarly, the CSub operation returns RLWE encryption of two ciphertext, here, as the description said, you need to replace the + sign with the - sign. The process remains the same.

2. Plaintext - Ciphertext Addition and Subtraction

This section discusses addition and subtraction between plaintexts and ciphertexts. This describes plaintext addition in detail. Just like the ciphertext-ciphertext version, the subtraction follows the same steps but replaces the + sign for a - sign. The plaintext-ciphertext addition operation is defined as $PPAdd(R, P)$, where R is an RLWE Ciphertext encrypting a polynomial $f(x)$ and P is an RLWE plaintext that encodes another polynomial $f_2(x)$. The operation $PPAdd$ returns a new RLWE ciphertext that encrypts $RLWE_{s(x)}(f(x) + p(x))$. Remember that an RLWE Ciphertext, $RLWE_{s(x)}(f(x)) = \langle a(x), b(x) \rangle$. Then, $PPAdd$ returns a new RLWE ciphertext $\langle a(x), b(x) + p(x) \rangle$.

3. Plaintext - Ciphertext Multiplication

The operation $PMul(R, P)$ multiplies an RLWE ciphertext $R = RLWE_{s(x)}(f(x))$ and an RLWE plaintext $p(x)$ and returns a new RLWE Ciphertext $RLWE_{s(x)}(f(x) \cdot p(x))$. Please note that while it says plaintext multiplication (which I'm also using to be consistent with the literature everywhere), the polynomial $p(x)$ isn't encoded, it's just the raw cleartext polynomial $p(x)$ which is multiplied. Recall that the RLWE ciphertext is $R = \langle a(x), b(x) \rangle$. The multiplication is done by multiplying both $a(x)$ and $b(x)$ by the polynomial $p(x)$. Note that this is a polynomial multiplication and not a pointwise multiplication and it's done mod $(x^N + 1)$. Therefore:

$$PMul(R, P) = \langle a(x) \cdot p(x), b(x) \cdot p(x) \rangle$$

Why did we just use cleartext $p(x)$ and not encode it

Remember that the ciphertext encrypts an encoded polynomial, whose coefficients are already multiplied by Δ . Encoding the plaintext polynomial and then multiplying will result in the resulting polynomial being 're-encoded', meaning, it becomes $(\Delta \cdot f(x)) \cdot (\Delta \cdot p(x)) = \Delta^2 \cdot f(x) \cdot p(x)$ which will mess up with the decryption. Therefore, we just need only one of the two encoded and since one of them is a ciphertext, its already encoded by design and thus, the other one should be left as cleartext. This is important to remember as literature uses plaintext multiplication and the terminology might give an incorrect impression.

```
# RLWECiphertext Code...
def add_ciphertext(self, other: RLWECiphertext):
    """
    Add another ciphertext <other> to the current ciphertext and return the result as a new
    ↪ ciphertext. Two add two ciphertexts, you just add their respective a and b
    ↪ polynomials.
    """
    # adding a and b, and creating a new ciphertext
    added_a = self.a.polynomial_add(other.a)
    added_b = self.b.polynomial_add(other.b)

    return RLWECiphertext(added_a, added_b, self.config)

def sub_ciphertext(self, other: RLWECiphertext):
    """
    Subtract another ciphertext <other> from the current one and return the result as a new
    ↪ ciphertext. This operation subtracts <other> from the <self> and not the other way
    ↪ round.
    """
    # subtracting a and b and returning a new ciphertext
    subbed_a = self.a.polynomial_subtract(other.a)
    subbed_b = self.b.polynomial_subtract(other.b)

    return RLWECiphertext(subbed_a, subbed_b, self.config)

def mul_plaintext(self, cx: RLWEPlaintext):
    """
    This function multiplies the given plaintext to self and returns the product as a new
    ↪ ciphertext.
    """
    # for this, you multiply both a and b with the plaintext
    mul_a = self.a.polynomial_multiply(cx.p)
    mul_b = self.b.polynomial_multiply(cx.p)

    return RLWECiphertext(mul_a, mul_b, self.config)
```

2.3.6 Trivial RLWE Ciphertext

A trivial RLWE ciphertext is an RLWE ciphertext where $a(x)$ is just a zero polynomial and $b(x)$ is just the encoded plaintext.

```
def get_trivial_rlwe_ciphertext(mx: Polynomial, config: RLWEConfig):
    """
    This function returns the trivial RLWE ciphertext for the given config. Trivial RLWE ciphertext
    ↪ is basically just an RLWE ciphertext where a is just a zero polynomial and b is polynomial
    ↪ itself.
    NOTE: Im not very sure if we should encode the input polynomial, meaning, if we take a cleartext
    ↪ input or plaintext input.
    """
```

```

'''
# sanity check
if len(mx.coeffs) != config.N:
    raise ValueError("The polynomial has a degree thats different from the degree provided in the
    ↪ config file. Please ensure that theyre the same.")
# creating a
a = make_zero_polynomial(config.N)
return RLWECiphertext(a = a, b = mx, config=config)

```

Chapter 3

GSW Ciphertexts and Ciphertext Multiplication

The previous section mentioned that multiplication between two RLWE ciphertexts is non-trivial. Instead, one way of multiplying two ciphertexts is multiplication between an RLWE and a GSW ciphertext. This chapter first describes GSW Ciphertext and then proceeds to talk about multiplying a GSW and an RLWE ciphertext together..

3.1 Base-p Representation

For ciphertexts, a polynomial's coefficients are encoded to the 'q-space', meaning that the coefficients will be in the range $[-\frac{q}{2}, \frac{q}{2})$. For the base-p representation, we define a parameter p (note that its not the same as the plaintext modulus p, anywhere in this section p will be used to denote p for base-p and the plaintext modulus will be pt), and Z_p denotes the range for elements to be in the range $[-\frac{p}{2}, \frac{p}{2})$. The base-p representation of a given integer $x \in Z_q$ will be a series of integers $(x_0, x_1, \dots, x_k)$ such that:

$$x = x_0 + x_1 \cdot p + x_2 \cdot p^2 + \dots + x_{k-1} \cdot p^{k-1}$$

For example, if we want to find the base-p representation of say $x = 256$ for $p = 2^8 = 256$ will be $[0, 1, 0, 0]$. For $x = 1000$, it will be $[-24, 4, 0, 0]$ (these are the vales of xis).

3.1.1 How are these calculated?

- To calculate the base-p representation of any integer, we need to find the length of the sequence. This is denoted by k and is calculated as $k = \frac{\log_2 q}{\log_2 p}$.
- After that, the easiest way to find the k integers for this representation is:
 - Get the binary representation of the integer. Even if the representation takes less than $\log_2 q$ bits, it should be expanded to fit the length.
 - Then, divide this representation into k chunks of length $\log_2 p$.
 - Each chunk converted into its integer representation will be the k-th element of the representation. Note that the sequence will start from the right hand side.

```
def get_k(q, p):  
    '''  
    Given q (ciphertext modulus) and p (not plaintext modulus, but a different parameter  
    → specifically for base p representation used for ciphertext-ciphertext multiplication of an  
    → RLWE and GWE cipherterats), get the value of k, which is the sequence length of base-p  
    → representation of any input.  
    '''  
    return np.int32(np.log2(q)/np.log2(p))
```

The problem with the above method to find the base- p representation is that it works with unsigned integers, as binary representation loses the sign and the elements of the representation will also be unsigned. But FHE works with signed integers. While this technique can't be used as-is, it can be used after making a small modification to the original integer. Basically, instead of finding the base- p representation of the integer x , it is offset by an integer B , and the above mentioned method is used to then calculate the base p representation of $x + B$. Let the representation be $(x'_0, x'_1, \dots, x'_{k-1})$. Once it is obtained, $\frac{p}{2}$ can be subtracted from the individual elements of the sequence to get the base- p representation of x .

How to find the offset value?

We know that due to the property of base- p representations,

$$x = (x'_0 - \frac{p}{2}) + (x'_1 - \frac{p}{2}) \cdot p + (x'_2 - \frac{p}{2}) \cdot p^2 + \dots + (x'_{k-1} - \frac{p}{2}) \cdot p^{k-1}$$

$$x = \sum_{i=0}^{k-1} (x'_i - \frac{p}{2}) \cdot p^i = \sum_{i=0}^{k-1} x'_i \cdot p^i - \frac{p}{2} \sum_{i=0}^{k-1} p^i$$

Rearranging this, we get,

$$x + \frac{p}{2} \cdot \sum_{i=0}^{k-1} p^i = \sum_{i=0}^{k-1} x'_i \cdot p^i$$

Based on this, $B = \frac{p}{2} \sum_{i=0}^{k-1} p^i$

```
def convert_array_to_base_p(a: np.ndarray, p, q):
    """
    This function converts each element into a its base p-representation, given the values of q
    ↪ (ciphertext modulus) and p (not plaintext modulus). Note that the representation for each
    ↪ element in a will be a sequence of a fixed length.
    """
    log_p = np.log2(p).astype(np.int32)
    k = get_k(q = q, p = p)
    p_half = p//2
    offset = p_half * np.sum([p**i for i in range(k)])
    mask = p - 1 # this in binary is exactly log_p ones.

    a_offset = np.uint32(a + offset) # need to set this to uint 32 as the the p-bit representation
    ↪ will be calculated from unsigned integers and then converted to signed
    # return a_offset

    output = [] # this will store the output base-p representation for each element
    # NOTE: output is a list of k numpy arrays. Each numpy array stores the ith p-bit representation
    ↪ of the corresponding original input.
    for i in range(k):
        output.append(
            (np.right_shift(a_offset, i * log_p) & mask).astype(np.int32) - np.int32(p_half)
        )

    return output

def convert_base_p_to_int_array(a: Sequence[np.array], p: np.int32):
    """
    This function takes as input the p-bit representations of an array, and converts it back into
    ↪ integer. a should be a list or a iterator of ndarrays where each ndarray contains the ith
    ↪ p-bit representation of the integer.
    """
    # calculate x from the representation: x = x_0 + x_1 p + ... + x_{k-1} p^{k-1}
```



```

decimal_repr = np.zeros(len(a[0]), dtype=np.int32)
for i in range(len(a)):
    decimal_repr += (a[i] * (p**i)).astype(np.int32)
return decimal_repr

# defining base p representation functions for polynomials
def get_base_p_from_polynomial(poly: Polynomial, p: np.int32, q: np.int32) -> Sequence[Polynomial]:
    """
    This function converts each coefficient of the coefficients of the given polynomial and returns
    ↪ them as a set of polynomials (sequence of the polynomials is very important).
    """
    coeffs = poly.coeffs
    # converting coeffs into base p representation
    coeffs_base_p = convert_array_to_base_p(coeffs, p, q)
    # create a polynomial for each array
    base_p_poly = [Polynomial(N=poly.N, coeffs = ic) for ic in coeffs_base_p]

    return base_p_poly

def get_poly_from_base_p(base_p_polys: Sequence[Polynomial], p) -> Polynomial:
    """
    This function calculates the polynomial from the given base p representation polynomials.
    """
    # get all the coefficients from the sequence
    base_p_coeffs = [p.coeffs for p in base_p_polys]
    # get the integer coeffs from the list
    int_coeffs = convert_base_p_to_int_array(base_p_coeffs, p)
    return Polynomial(N=base_p_polys[0].N, coeffs = int_coeffs)

```

Why is the Base-p representation needed?

The base-p representation will be needed to perform the GSW encryption down the line which will be used for multiplication.

3.2 GSW Encryption

The idea behind a GSW encryption is to 'hide/mask' a plaintext polynomial $f(x)$ between **zero ciphers** to encrypt it and hide it.

```

class GSWConfig:
    """
    This class implements the GSW Config. It takes the RLWE config, as most of the configurations
    ↪ are identical. Additionally, it also takes gsw_p, which will be used for the base p
    ↪ representation.
    """
    def __init__(self, rlwecfg: RLWEConfig, gsw_p: int):
        self.rlwecfg = rlwecfg
        self.gsw_p = gsw_p

```

- **Zero Cipher (RLWE):** A zero cipher is just the RLWE encryption of 0. Note that it is an actual encryption where the input cleartext (and thus the resulting plaintext) are 0, however, it still uses the entire encryption mechanisms and thus, the polynomials $a(x)$ and $b(x)$ are not zero. Remember this detail because in my experience, it becomes easy to mistake it for something that's more like a trivial RLWE encryption otherwise.

Let $Z_0, Z_1 \in RLWE_{s(x)}(0)$ be two RLWE zero ciphers, and Z be a 2×2 matrix whose i -th row is Z_i (the matrix has 2 columns if you count $a(x)$ and $b(x)$ as separate columns, which is what's being done here.)

$$Z = \begin{pmatrix} Z_1 \\ Z_2 \end{pmatrix}$$

And let the 2 x 2 identity matrix $I_{2 \times 2} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$

The GSW ciphertext of a given polynomial is then given by:

$$Enc_{s(x)}^{GSW}(f(x)) = f(x) \cdot I_{2 \times 2} + Z$$

```
class GSWConfig:
    """
    This class implements the GSW Config. It takes the RLWE config, as most of the configurations
    → are identical. Additionally, it also takes gsw_p, which will be used for the base p
    → representation.
    """
    def __init__(self, rlwecfg: RLWEConfig, gsw_p: int):
        self.rlwecfg = rlwecfg
        self.gsw_p = gsw_p
```

```
class GSWCiphertext:
    """
    This class implements the gsw ciphertext.
    """
    def __init__(self, config: GSWConfig, polys: Sequence[RLWEciphertext]):
        self.config = config
        self.polys = polys
```

```
class GSWEncryptionKey:
    """
    This class implements the GSW encryption key and methods to encrypt and (possibly) decrypt a gsw
    → ciphertext.
    """
    def __init__(self, config: GSWConfig):
        self.config = config
        self.key = Polynomial # this is just going to be the RLWE encryption key s(x)
        self.rlwe_key = RLWEEncryptionKey

    def generate_key_from_rlwe_key(self, rlwe_key: RLWEEncryptionKey):
        """
        This function creates the GSWEncryptionKey from the RLWEEncryption key. To be honest,
        → its just the RLWEEncryption key as it is, but for the sake of structure of code and
        → to keep it consistent with other classes. Its a little silly to save both the
        → rlwe_key object and the key polynomail as well, but eh...
        """
        self.key = rlwe_key.key
        self.rlwe_key = rlwe_key

    def get_rlwe_key_from_gsw(self):
        """
        This function returns the rlwe key from the current gsw_key. Again, its the same, but
        → consistency, formality, something, something...
        """
        rlwe_key = RLWEEncryptionKey(config=self.config.rlwecfg)
        rlwe_key.key = self.key
        return rlwe_key
```

```

def encrypt(self, fx: Polynomial):
    """
    Encrypt a given polynomial as a GSWCiphertext.
    """
    # first, we need to get the value of k
    k = get_k(self.config.rwlecfg.q, self.config.gsw_p)
    gsw_p = np.int32(self.config.gsw_p) # i dont want to write the whole thing
    # make Z: the zero encryption array of shape (2k, 2), note that each row is a RLWE
    # → ciphertext, the 2 just means (a(x), b(x))
    zero_poly = make_zero_polynomial(self.config.rwlecfg.N)
    zero_poly_plain = encode_rlwe(zero_poly, self.config.rwlecfg) # while for 0 config, the
    # → coefficients for both the plaintext and ciphertext are going to be the same, the
    # → encryption for RLWE only accepts an RLWEPlaintext and not a Polynomial
    Z = [
        self.rlwe_key.encrypt(zero_poly_plain) for _ in range(2*k)
    ] # this has shape (2k, 2)
    # create the GSW ciphertext
    # its created using the equation: fx * B_p + Z
    # i will implicitly divide the Z matrix in two halves
    # here, ill iterate for k times, and for each iteration, the corresponding index's RLWE
    # → of the first half's a gets fx * (p**i) added and the corresponding index of the
    # → second half's b gets the same added
    for i in range(k):
        # multiply fx with the correct power of p
        scaled_fx = fx.polynomial_const_multiply(gsw_p**i)
        # add it to first half's a
        Z[i].a = Z[i].a.polynomial_add(scaled_fx)
        # add it to second half's b
        Z[i + k].b = Z[i + k].b.polynomial_add(scaled_fx)

    return GSWCiphertext(self.config, Z)

```

This results in a **vector** of two RLWE ciphertexts (note that it says vector of 2 RLWE ciphertexts, while Z is denoted as a matrix, each row there is an RLWE ciphertext, so it is also a vector to 2 RLWE ciphertexts. The reason for this weird seeming notation will be clear later.). There are a few things to be noted about GSW ciphertext/encryption:

- The multiplication between $f(x)$ and the identity matrix is pointwise.
- GSW ciphertext doesn't have to be a vector of just two RLWE ciphertexts, it can be arbitrarily long, just that each element of this vector has to be an RLWE ciphertext. In the coming section(s), a much longer GSW ciphertext will be created.

3.3 RLWE - GSW Multiplication

3.3.1 CMul: Basic Idea

If you have two polynomials $g(x)$ and $f(x)$, with an RLWE ciphertext $R = \text{Enc}_{s(x)}^{\text{RLWE}}(f(x))$ and a GSW ciphertext $G = \text{Enc}_{s(x)}^{\text{GSW}}(g(x))$, there is a function CMul such that:

$$R_{\text{prod}} = \text{CMul}(R, G) = G \cdot R$$

And CMul returns an RLWE ciphertext that encrypts $f(x) \cdot g(x)$. Looking further into CMul:

$$R_{\text{prod}} = R \cdot (g(x) \cdot I_{2 \times 2} + Z)$$

This method will work as long as the coefficients of $a(x)$, $b(x)$ and $g(x)$ are small. Based on the definition of GSW encryption, the polynomial $g(x)$ will have (relatively) small coefficients (as it's not really being encoded or encrypted),

however, since $f(x)$ is being encrypted into RLWE ciphertexts, coefficients of $a(x)$ and $b(x)$ will be very large. To decrease these magnitudes, one of the simplest ways can be by just multiplying it with $\frac{1}{p}$, where $p \nmid 1$, this will scale the coefficients of both $a(x)$ and $b(x)$ down and satisfy the condition. To compensate for the the GSW ciphertext can be scaled up and defined like so:

$$G = p \cdot f(x)I_{2 \times 2} + Z$$

Then, CMul becomes:

$$CMul(R, G) = \frac{1}{p} \cdot R \cdot G = \frac{1}{p} \cdot R(p \cdot f(x)I_{2 \times 2} + Z)$$

p and $\frac{1}{p}$ cancel each other out and we get the same result, but this operation ensures that all the polynomials have small coefficients. However, there is a problem with this approach. For scaling the coefficients down, division by p is done. However, due to the way integers in the ring R_q are, division by p is not well defined. It can work if p is set so that q is divisible by p , however, that leads to a loss of precision, which is not the best idea.

3.3.2 CMul with base p representation

Instead of scaling down an integer by dividing it with p , its base- representation can also be used.

$$x = x_0 + x_1 \cdot p + \dots + x_{k-1}p^{k-1}$$

If coefficients of polynomials are decomposed using this, it also ensures that the resulting decomposed coefficients remain small (as long as p is chosen to be relatively small, which it often is) but it avoids the loss of precision unlike the previous approach. There is one tradeoff however, if a polynomial's coefficients are decomposed using the base p representation, it results in k polynomials for a single one, which can increase runtimes whenever any operation uses this decomposition method. The base p representation of a polynomial $f(x)$ is a set of polynomials $[f_0(x), f_1(x), \dots, f_{k-1}(x)]$ such that:

$$f(x) = f_0(x) + p \cdot f_1(x) + \dots + p^{k-1} \cdot f_{k-1}(x)$$

Where each polynomial $f_i(x)$'s coefficients are the i th base p representation of the corresponding polynomial in $f(x)$. The above equation can also be written as a dot product:

$$f(x) = (f_0(x), f_1(x), \dots, f_{k-1}(x)) \cdot (p^0, p^1, \dots, p^{k-1})$$

Then the polynomial $f(x)$ can be defined as:

$$f(x) = Base_p(f(x)) \cdot v_p$$

where $v_p = (p^0, p^1, \dots, p^{k-1})$ (This is a row vector.)

Using this logic, the components $a(x)$ and $b(x)$ for an RLWE ciphertext R can be decomposed into their base p representations (remember that this way we reduce the magnitude of their coefficients) such that:

$$R = (Base_p(a(x)), Base_p(b(x))) \cdot B_p$$

Where B_p is a matrix of shape $2k \times 2$, $B_p = \begin{pmatrix} v_p^t & 0_p \\ 0_p & v_p^t \end{pmatrix}$, 0_p being a column vector of zeros of size k , and $(Base_p(a(x)), Base_p(b(x)))$ is the concatenation of base p representations of $a(x)$ and $b(x)$.

Here, the GSW encryption also changes slightly. It is defined as:

$$G = Enc_{s(x)}^{GSW}(g(x)) = g(x) \cdot B_p + Z$$

Where B_p the same as defined above, but Z is a $(2k \times 2)$ matrix, where each row is RLWE encryption of 0 and thus the GSW encryption becomes a $(2k \times 2)$ matrix where each row is an RLWE ciphertext. The CMul now changes to:

$$CMul(R, P) = (Base_p(a(x)), Base_p(b(x))) @ G$$

This multiplication is sort of like a matrix-multiplication between the two matrices. Remember that $(Base_p(a(x)), Base_p(b(x)))$ is concatenated base p representations, which makes a vector of $2k$ polynomials, and multiplying it with G , which is a $(2k \times 2)$ matrix (each row is an RLWE ciphertext $(a(x), b(x))$), and a matmul between then results in a vector of two polynomials, which will be the corresponding $\langle a(x), b(x) \rangle$ RLWE ciphertext that encrypts the ciphertext multiplication of two polynomials.

```
def ciphertext_multiplication(G: GSWCiphertext, R: RLWECiphertext) -> RLWECiphertext:
    """
    This function performs the ciphertext-ciphertext multiplication between G and R. G is of shape
    → (2k, 2), basically 2k RLWE Ciphertexts and each Rlwe ciphertext is (a, b).
    So a little explanation because this function looks a little confusing (but in reality it really
    → isnt once you look at it hard enough keeping in mind how the multiplication is actually
    → done). Im going to define in the docstring how does this operation actually works.
    So:
    First, G is a series of 2k RLWECiphertexts. Essentially, theyre all zero ciphertexts to which fx
    → (the polynomial with which we actually want to achieve multiplication with R) is added. For
    → the first half of G [0:k-1], fx is added to a(x) of each RLWE 'row', and is added to b(x)
    → for the second half [k: 2k-1].
    Second, a(x) and b(x) of R are converted to corresponding base-p representation. Now R was
    → basically a vector/sequence/whatever of 2 polynomials, but now it becomes a vector of 2k
    → polynomials (k polynomials for base p representation of a(x) and b(x)).
    To perform c-c multiplication, you basically do something akin to matrix multiplication,
    → (Basep(a(x)), Basep(b(x))) @ G, which gives you a single rlwe polynomial, and its shape will
    → be (1, 2), a single (a(x), b(x)) ciphertext. Here, each multiplication is polynomial
    → multiplication followed by polynomial addition for matmul.
    """
    gsw_config = G.config
    rlwe_config = R.config
    # first get the base p representation for R, for both a(x) and b(x) and create it into a
    → sequence of 2k polynomials
    R_base_p_repr = get_base_p_from_polynomial(R.a, p=gsw_config.gsw_p, q=rlwe_config.q) +
    → get_base_p_from_polynomial(R.b, p=gsw_config.gsw_p, q=rlwe_config.q) # the output of the
    → function is a list so + just appends

    # so matmul here is implemented using a for loop, first, an RLWECiphertext that is basically all
    → zeros is initialized and then the corresponding products a.G and b.G for each R_base_p_repr
    → is added to a and be of this new RLWECiphertext respectively
    res_rlwe_cipher = RLWECiphertext(
        config = rlwe_config,
        a = make_zero_polynomial(rlwe_config.N),
        b = make_zero_polynomial(rlwe_config.N)
    )

    for i, rp in enumerate(R_base_p_repr):
        # multiply the current r to both a and b of the current row of g.a
        a_mul = rp.polynomial_multiply(G.polys[i].a)
        # adding
        res_rlwe_cipher.a = res_rlwe_cipher.a.polynomial_add(a_mul)

        # doing the same for b
        b_mul = rp.polynomial_multiply(G.polys[i].b)
        # adding
        res_rlwe_cipher.b = res_rlwe_cipher.b.polynomial_add(b_mul)

    return res_rlwe_cipher
```

3.4 GSW Encryption Key

The actual encrypted ciphertexts in a GSW ciphertexts are the zero RLWE ciphertexts, we do not actually encrypt the polynomial $f(x)$. And since these are RLWE ciphertexts, the gsw encryption key polynomial is the same as the rlwe encryption key polynomial.

Chapter 4

Some More Operations: BlindRotate & SampleExtract

These are operations that are used in bootstrapping.

4.1 The Homomorphic Multiplexer

4.1.1 What is a multiplexer gate?

A multiplexer is a gate/function that takes two inputs l_0, l_1 and a selector bit b . The output of a multiplexer is defined as:

$$Mux(b, l_0, l_1) = \begin{cases} l_0 & \text{if } b = 0 \\ l_1 & \text{if } b = 1 \end{cases} \quad (4.1)$$

Basically, it outputs one of the two inputs based on the selector bit. The multiplexer operation can also be defined as:

$$Mux(b, l_0, l_1) = b(l_1 - l_0) + l_0$$

How to make this with with ciphertexts

In a homomorphic multiplexer, the selector bit is an GSW Ciphertext, while the two inputs are RLWE Ciphertexts. Let B be the GSW encryption of the selector bit and L_0 and L_1 be the RLWE ciphertext inputs. Then, the CMux operation is defined as:

$$CMux(B, L_0, L_1) = CAdd(CMul(B, CSub(L_1, L_0)), L_0)$$

The respective methods are already defined and can be used to implement this CMux operation using the relevant ciphertexts. Remember that the keys need to be from the same origin (the RLWE key should be converted into the GSW key for encrypting the selector).

```
def cmux(b: GSWCiphertext, l0: RLWECiphertext, l1: RLWECiphertext) -> RLWECiphertext:
    """
    This function implements the multiplexer for ciphertexts.
    The homomorphic multiplexer operation homomorphically computes: b(l1 - l0) + l0.
    So its basically chaining the following operations: CAdd(CMul(b, CSub(l1, l0)), l0)
    the function returns l0 is b == 0, else l1.
    """
    # calculating l1 - l0
    subbed = l1.sub_ciphertext(l0)
```

```

# calculating product of subbed and b
prod = ciphertext_multiplication(b, subbed)
# adding 10 to get the cmux result
mux_cipher = prod.add_ciphertext(10)

return mux_cipher

```

4.2 BlindRotate

Take a polynomial $f(x) = 1 + 2x + x^2 \in \mathbb{Z}_q[x]/(x^3 + 1)$. Now do:

$$f(x) \cdot x = x + 2x^2 + x^3$$

which is equal to:

$$x \cdot f(x) = -1 + x + 2x^2 \pmod{(x^3 + 1)}$$

If you stare at the coefficients, you can see that it 'rotated' with an opposite sign, it went from $[1, 2, 1]$ to $[-1, 1, 2]$. This is why, multiplication by a monomial x^i is also called rotating a polynomial. The homomorphic version of this rotation is called BlindRotate. First we will describe the cleartext version of this operation and then move to the HE version. First, lets take a look at rotating a polynomial with an integer i . This is denoted as:

$$\text{IntRotate}(i, f(x)) = x^i \cdot f(x) \text{ [This is the cleartext version]}$$

Note that the i in x^i can be both positive and negative, but it has to be an integer. Also, due to the negacyclic property of the polynomials in the ring,

$$\text{IntRotate}(i + 2N, f(x)) = \text{IntRotate}(i, f(x))$$

This is because the polynomials in the ring have a period of $2N$. However, there is a problem: while the polynomial has a period of $2N$, the coefficients do not, as they are R_q . Therefore, to get this to work properly with ciphertexts, we have to scale i to be of the same scale. This can easily be done by rescaling $i \in \mathbb{Z}_q$ by multiplying it with $\frac{2N}{q}$ and then using the rescaled version to rotate $f(x)$. This operation is called Rotate:

$$\text{Rotate}(i, f(x)) = \text{IntRotate}(\lfloor \frac{2N}{q} \cdot i \rfloor, f(x)) = x^{\lfloor \frac{2N}{q} \cdot i \rfloor} \cdot f(x)$$

For example, if you want to rotate $f(x)$ by $\frac{N}{4}$, you need to set $i = \frac{q}{8}$.

Now, coming to the homomorphic version of the operation. It is called BlindRotate and it takes as input an LWE ciphertext encrypting i (I), and an RLWE ciphertext encrypting $f(x)$ (R). It is written as $\text{BlindRotate}(I, R) = \text{RLWE}_{s(x)}(x^{\lfloor \frac{2N}{q} \cdot i \rfloor} \cdot f(x))$.

(This will be relevant in the future when discussing TFHE specifically). The scaling of i in BlindRotate, to change it from $\mathbb{Z}_q$ to $\mathbb{Z}_{2N}$ is called modulus switching.

How is this actually implemented using ciphertexts?

For easier notation, lets have $\alpha = \frac{2N}{q}$. $I = \text{LWE}_s(i) = \langle a, b \rangle$ and $\text{RLWE}_{s(x)}(f(x)) = \langle a(x), b(x) \rangle$ are the inputs to the BlindRotate function. First, lets define a' and b' as:

$$\begin{aligned} a' &= \lfloor \alpha \cdot a(x) \rfloor \ \& \\ b' &= \lfloor \alpha \cdot b(x) \rfloor \end{aligned}$$

$\langle a', b' \rangle$ is the LWE encryption of $\text{PMul}(\alpha, I)$. Also, since decryption of an LWE ciphertext is given as $i + e = b - a \cdot s$, and since $\alpha \cdot I$ is a plaintext-ciphertext multiplication between a constant and an LWE ciphertext, the decryption will be $b' - a' \cdot s$.

Thus,

$$x^{\lfloor \alpha(i+e) \rfloor} \cdot f(x) = x^{b'-s \cdot a'} \cdot f(x) = x^{b'} \prod_{i=1}^N (x^{a'_i \cdot s_i}) \cdot f(x)$$

Note that while ideally, i is from 0 to $N-1$, but here, its from 1 to N . There is a reason for that. Also, here, s_i is binary (either 0 or 1), and depending on its value, $a'_i \cdot s_i$ can either be 0 or a'_i . Additionally, if you define $g_0(x) = x^{b'} \cdot f(x)$, then the above equation can be modeled iteratively such that:

$$g_i(x) = g_{i-1}(x) \cdot x^{a'_i \cdot s_i} \text{ [Think about it, do some examples manually and it be clear why]}$$

Combine this with the fact that the s_i is either a 0 or a 1, this can be modeled using a CMux operation with s_i being the selector bit. For $i = 1, 2, \dots, N$:

$$g_i(x) = \begin{cases} g_{i-1}(x) & \text{if } s_i = 0 \\ g_{i-1}(x) \cdot x^{a'_i} & \text{if } s_i = 1 \end{cases} \quad (4.2)$$

However, to be able to do this, we need to make a change/do an additional step. Note that for CMux, the selector bit needs to be a GSW ciphertext, however, the secret key array isnt that. Solution, encrypt each component of the secret key s_i into a GSW ciphertext, $t_i = GSW_{s(x)} s_i$. The resulting vector of GSW ciphertexts $[t_0, t_1, \dots, t_n]$ is called the **Bootstrapping Key**.

```
class BootstrapKey:
    """
    This class implements the bootstrap keys and the methods associated with it, like BlindRotate.
    """
    def __init__(self, config: GSWConfig):
        self.config = config
        self.key = Sequence[GSWCiphertext] # each component (s_i) from the LWE key will be a GSW
        ↪ ciphertext

    def generate_bootstrap_key(self, lwe_key: LWEEncryptionKey, gsw_key: GSWEncryptionKey):
        """
        generate the bootstrapping keys. The bootstrapping keys are just GSW encryption of each
        ↪ component of the LWE_key using the gsw_key provided.
        NOTE TO SELF: Im debating whether to make the lwe_key and gsw_key members of this class
        ↪ or not. Right now I wont, but if there are other operations that require these keys
        ↪ in addition to the bootstrap keys, I will.
        """
        bs_keys = [] # this will store the keys
        # looping across each lwe key component and encrypting it
        for si in lwe_key.key:
            # first, create a monomial with the si being constant
            si_monomial = build_monomial(self.config.rwlecfg.N, 0, si)
            si_gsw = gsw_key.encrypt(si_monomial)
            bs_keys.append(si_gsw)
        self.key = bs_keys
```

Other than this, we are all set. Now, the Blindrotate operation can be defined using a combination of the methods that we already have defined. First, the relation

$$G_i = CMux(t_i, G_{i-1}, PMul(x^{a'_i}, G_{i-1})),$$

where, $G_0 = PMul(x^{b'}, F)$
(F is the RLWE ciphertext encrypting $f(x)$)

Then, the BlindRotate operation is defined as:

$$BlindRotate(I, F) = G_N$$

In case you forgot where I is actually used (I will sometimes), due to the wall of text and a lot of math above, remember that the scaled version of I , $\alpha \cdot I$, is $\langle a', b' \rangle$.

```
# BootstrapKey Class...
def blindrotate(self, FX: RLWECiphertext, I: LWECiphertext) -> RLWECiphertext:
    """
    This function implements the Rotate(I, FX) function homomorphically. This function,
    ↪ given i, and a polynomial f(x) (I = LWE(i), FX = RLWE(f(x)), calculates  $x^i \cdot f(x)$ 
    ↪ as an RLWECiphertext.
    """
    # scale the lwe ciphertext
    scaled_I_a = np.int32(np rint(np.multiply(I.a, ((2 * self.config.rwlecfg.N)/
    ↪ self.config.rwlecfg.q))))
    scaled_I_b = np.int32(np rint(np.multiply(I.b, ((2 * self.config.rwlecfg.N)/
    ↪ self.config.rwlecfg.q)))) # NOTE: This is going to be a constant since this is an
    ↪ LWE ciphertext

    # rotate operation
    # initialize g0
    xb_poly = build_monomial(self.config.rwlecfg.N, scaled_I_b, 1)
    xb_plain = RLWEPlaintext(xb_poly, FX.config)
    rotated_poly = FX.mul_plaintext(xb_plain) # this will updated every timestep
    # loop N times to finally get the rotated output
    for i, ai in enumerate(scaled_I_a):
        # get  $g_{i-1}(x) \cdot x^{a_i}$ 
        poly = RLWEPlaintext(build_monomial(self.config.rwlecfg.N, -ai, 1),
        ↪ self.config.rwlecfg) # just converting it to a plaintext because the
        ↪ plaintext multiply in RLWE wants a plaintext
        right = rotated_poly.mul_plaintext(poly)
        # update rotated_poly using the cmux operation
        rotated_poly = cmux(self.key[i], rotated_poly, right)

    return rotated_poly
```

4.3 Extract Sample

The SampleExtract function homomorphically returns the i -th coefficient of the encrypted RLWE polynomial as an LWE Ciphertext. The returned LWE encryption of the coefficient is encrypted under the associated LWE encryption key of the RLWE encryption key. Note that for any RLWE ciphertext, this is true:

$$f(x) + e(x) = b(x) - a(x)s(x) \text{ (f(x) is the scaled plaintext)} \quad \text{--- (i)}$$

Let $(a(x) \cdot s(x))_i$ be the i -th coefficient of $a(x) \cdot s(x)$. Due to the negacyclic property of these functions, the i -th coefficient satisfies the following polynomial convolution property:

$$(a(x) \cdot s(x))_i = \sum_{j=0}^i a_{i-j} s_j - \sum_{j=i+1}^{N-1} a_{N+i-j} s_j$$

The $Conv(a(x), i)$ vector can be defined as:

$$Conv(a(x), i) = (a_i, a_{i-1}, \dots, a_1, a_0, -a_{N-1}, -a_{N-2}, \dots, -a_{i+1}) \text{ (each } a_i \text{ is the } i\text{th coefficient of } a(x))$$

Then $(a(x) \cdot s(x))_i$ can be defined as a dot product: $Conv(a(x), i) \cdot s$. Also, from equation (i), we can specifically get the i -th coefficient like this:

$$f_i + e_i = b_i - (a(x) \cdot s(x))_i$$

Replacing this with the new expression for $(a(x), s(x))_i$:

$$f_i + e_i = b_i - \text{Conv}(a(x), i) \cdot s$$

Based on this, the LWE ciphertext for the extracted coefficient becomes: $\text{ExtractSample}(i, F) = (\text{Conv}(a(x), i), b_i)$

```
def extract_sample(i: int, FX: RLWECiphertext) -> LWECiphertext:
    """
    This function extracts the ith coefficient of the polynomial encrypted in the input FX and
    ↪ returns the ith coefficient as an LWE Polynomial. It calculates the ai (a for the LWE
    ↪ ciphertext) using Conv(a(x), i) and the ith coefficient of b(x). The resulting LWE
    ↪ ciphertext is encrypted using the LWE key obtained from the RLWE key (you can convert them
    ↪ into one another).
    """
    # first get the conv(a(x), i)
    # conv(a(x), i) is just going to be coeffs of a(x), but 'rotated' around x.
    # that means, if a(x) was [a0, a1, a2, ..., a(N-1)]
    # then conv(a(x), i) is [ai, a_{i-1}, ..., a1, a0, -a_{N-1}, -a_{N-2}, ..., -a_{i+1}]
    ax = FX.a.coeffs
    conv_ax = np.hstack([
        ax[:i+1][::-1],
        -1 * ax[i+1:][::-1]])
    bi = FX.b.coeffs[i] # just need the ith coefficient's b

    # creating an LWECiphertext for this
    return LWECiphertext(
        a = conv_ax,
        b = bi
    )
```

Chapter 5

Bootstrapping

This section describes how bootstrapping works in FHE, by trying to implement a NAND operation between two ciphertexts. But before starting with bootstrapping, little bit of a context.

5.1 NAND and Homomorphic NAND

NAND is a boolean operation that takes two inputs and outputs one depending on the inputs. For two boolean inputs p & q , $\text{NAND}(p, q)$ is defined as:

$$\text{NAND}(p, q) =$$

p	q	$\text{NAND}(p, q)$
T	T	F
T	F	T
F	T	T
F	F	T

The homomorphic NAND will take as input two LWE Ciphertexts, and return an output which will be the LWE encryption of the NAND operation done on the values the input ciphertexts encrypt. Mathematically, for two ciphertexts $L_0 = \text{LWE}(l_0)$ and $L_1 = \text{LWE}(l_1)$:

$$\text{CNAND}(L_0, L_1) = \text{LWE}(\text{NAND}(l_0, l_1))$$

The reason for choosing this operation to explain bootstrapping is that its a fairly common operation that can be seen in many FHE circuits, where it is also chained multiple times. However, in arithmetic operations defined above, the noise in the ciphertext increases, which is not good if youre chaining these operations. Therefore, the methodology used to implement CNAND should be such that it refreshes the noise, meaning, the noise of the result should not be a function of the noise of the input ciphertexts. This is why, bootstrapping is used. The way bootstrapping is defined ensures that the noise in the output does not blow up and as a result, bootstrapping operations can be chained together.

5.2 Encrypting Booleans

Booleans are represented as integers which are then encrypted. True is represented by 2 and False is represented by 0.

5.3 Implementing NAND

For this chapter, we go with the following parameter values:

- Plaintext modulus (p): 8 $[-4, 4)$
- Ciphertext modulus (q): $2^{32} = [-2^{31}, 2^{31})$

So, $\text{Encode}(i) = i \cdot 2^{29}$. Also, $\text{True} = \text{Encode}(2)$ and $\text{False} = \text{Encode}(0)$.




Figure 5.1: Ring for the above mentioned parameters, image taken from the blog.

Consider the function $F(b_0, b_1)$, where both inputs are LWE ciphertexts:

$$F(b_0, b_1) = \text{Encode}(-3) - b_0 - b_1.$$

Looking at the function and the ring above, the function starts at $\text{Encode}(3) = -3 \cdot 2^{29}$ and essentially just moves two dots anti-clockwise for each $b_i = \text{Encrypt}(2)$ (meaning, it is true).

$$F(b_0, b_1) =$$

b_0	b_1	$F(b_0, b_1)$
$\text{Encode}(2)$	$\text{Encode}(2)$	$\text{Encode}(1)$
$\text{Encode}(2)$	$\text{Encode}(0)$	$\text{Encode}(3)$
$\text{Encode}(0)$	$\text{Encode}(2)$	$\text{Encode}(3)$
$\text{Encode}(0)$	$\text{Encode}(0)$	$\text{Encode}(-3)$

Note that when both b_0 and b_1 are $\text{Encode}(2)$, $-2 \cdot 2^{29} < F(b_0, b_1) < 2 \cdot 2^{29}$. We can define a step function such that:

$$\text{Step}(x) = \begin{cases} 0 & \text{if } -q/4 < x \leq q/4 = -2 \cdot 2^{29} < x \leq 2 \cdot 2^{29} \\ \text{Encode}(2) & \text{otherwise} \end{cases} \quad (5.1)$$

Then, $\text{Step}(F(b_0, b_1))$ is basically the NAND operation.

$$\text{Step}(F(b_0, b_1)) =$$

b_0	b_1	$F(b_0, b_1)$	$\text{Step}(F(b_0, b_1))$
$\text{Encode}(2)$	$\text{Encode}(2)$	$\text{Encode}(1)$	$\text{Encode}(0)$
$\text{Encode}(2)$	$\text{Encode}(0)$	$\text{Encode}(3)$	$\text{Encode}(2)$
$\text{Encode}(0)$	$\text{Encode}(2)$	$\text{Encode}(3)$	$\text{Encode}(2)$
$\text{Encode}(0)$	$\text{Encode}(0)$	$\text{Encode}(-3)$	$\text{Encode}(2)$

Therefore, we can say that $\text{NAND}(l_0, l_1) = \text{Step}(F(b_0, b_1)) = \text{Step}(\text{Encode}(-3) - b_0 - b_1)$.

The subtraction methods defined above can be used to calculate the value of F , and bootstrap can be used to calculate the Step function. Thus, this operation becomes equivalent to $\text{Bootstrap}(\text{CSub}(\text{Encode}(-3), \text{CSub}(b_0, b_1)))$. Here, -3 becomes a trivial LWE ciphertext instead of plaintext.

5.4 Bootstrap: Building the Step Function

Consider a polynomial $t(x) = -1 - x - x^2 - \dots - x^{N/2-1} + x^{N/2} + x^{N/2+1} + \dots + x^{N-1}$, where $N/2$ coefficients are -1 and the other half are 1 . We call this the test polynomial. Now, if this is rotated by 1 :

$$\text{IntRotate}(1, t(x)) = x \cdot t(x) = -1 - x - \dots - x^{N/2} + x^{N/2+1} + \dots + x^{N/2}$$

```
def generate_tx(N) -> Polynomial:
    """
    This function generates the polynomial i(t).
    """
    # building the polynomial
    tx_coeffs = np.ones(N, dtype=np.int32)
    # updating the first half to be -1
    tx_coeffs[:N//2] = -1
    tx_poly = Polynomial(N, tx_coeffs)
    return tx_poly
```

This rotates the coefficients one place to the right. The coefficient of x^{N-1} comes to the front with a - sign. Meaning, now, there are $N/2 + 1$ coefficients that are -1 and $N/2 - 1$ that are 1. Now, if we keep rotating it, such that we rotate it by $N/2$:

$$\text{IntRotate}(N/2, t(x)) = -1 - x - \dots - x^{N-1}$$

All the coefficients are -1. Rotating it one more time would result in a polynomial:

$$\text{IntRotate}(N/2 + 1, t(x)) = 1 - x - x^2 - \dots - x^{N-1}$$

Rotating it N times results in all the coefficients becoming positive. So, if we were to try and just look at the coefficient of the constant term (x^0), we can generalize it as follows:

$$\text{Coeff}(0, \text{IntRotate}(i, t(x))) = \begin{cases} -1 & \text{if } 0 \leq i \leq N/2 \\ 1 & \text{if } N/2 < i \leq N \end{cases} \quad (5.2)$$

This equation can be further generalized to include cases when i is negative, which is also possible

$$\text{Coeff}(0, \text{IntRotate}(i, t(x))) = \begin{cases} -1 & \text{if } -N/2 \leq i \leq N/2 \\ 1 & \text{otherwise} \end{cases} \quad (5.3)$$

Also, from the blind rotate section, we know:

$$\text{Rotate}(i, f(x)) = \text{IntRotate}(\lfloor \frac{2N}{q} \cdot i \rfloor, f(x))$$

So we can replace the `IntRotate` function with the `Rotate` function (because a. its defined for ciphertexts, and b. it will the change the limits from N to q , which is what we have when working with ciphertexts):

$$\text{Coeff}(0, \text{Rotate}(i, t(x))) = \begin{cases} -1 & \text{if } -q/4 \leq i \leq q/4 \\ 1 & \text{otherwise} \end{cases} \quad (5.4)$$

The above equation looks very similar to the step function, it has the same conditions for the input, but the output differs as this outputs -1 or 1 and step function outputs `Encode(2)` or 0. We can modify the above function to get the output of Step like this:

$$\text{Step}(i) = \frac{\text{Encode}(2)}{2} + \text{Coeff}(0, \text{Rotate}(i, \frac{\text{Encode}(2)}{2} \cdot t(x)))$$

Remember that `Encode(2)` is not a plaintext but a trivial LWE encryption of 2. This way, we have the step function implemented as bootstrap, and it utilizes `BlindRotate` and `Coeff` (`ExtractSample`) functions to implement it homomorphically.

```

def create_trivial_rlwe_ciphertext(mx: Polynomial, config: RLWEConfig):
    """
    This function creates a trivial RLWE ciphertext. A trivial RLWE ciphertext is one where the
    ↪ polynomial is one where b is the polynomial itself while a is a zero polynomial.
    """
    a = make_zero_polynomial(config.N)
    return RLWECiphertext(a, mx, config)

def create_trivial_lwe_ciphertext(m: LWEPlaintext, config: LWEConfig):
    """
    This function creates a trivial LWE ciphertext, which is an LWECiphertext with a being a vector
    ↪ of zeros and b being the plaintext itself.
    """
    a = np.zeros(config.n, dtype=np.int32)
    return LWECiphertext(a, m.m)

```

```

def bootstrap(i: LWECiphertext, bsk: BootstrapKey, R: RLWECiphertext, scale_lwe: LWECiphertext):
    """
    This function implements the bootstrap function to calculate the step operation defined above
    ↪ for the input i. The bsk is the bootstrapping key.
    This function returns 0 if i is in  $(-q/4, q/4]$ , else, it returns an lwe encryption of scale.
    NOTE [IMPORTANT]: Note that if scale is the cleartext np.int32 value, scale_lwe encrypts
    ↪ Encode(scale)/scale.
    """
    # blind rotate R
    rotated_rlwe = bsk.blindrotate(R, i)
    # extract sample to get the coefficient of constant
    coeff_lwe = extract_sample(0, rotated_rlwe)
    # adding scale_lwe to the output to get the offset-ed result of bootstrapping
    result_lwe = scale_lwe.add_ciphertext(coeff_lwe)

    return result_lwe

```

5.4.1 Details for FHE Implementation of Step using Bootstrapping

This section will describe the implementation details for the above defined operation. We want to implement bootstrap that evaluates the step function. Given an LWE ciphertext L (which will be the output of F described in the previous sections). The following things are needed:

- $L \in \text{LWE}_s(i)$, an LWE encryption of i .
- $R \in \text{RLWE}_{s(x)}(\frac{\text{Encode}(2)}{2} \cdot t(x))$, a trivial RLWE encryption of the scaled test polynomial.
- $B \in \text{LWE}_s(\frac{\text{Encode}(2)}{2})$, a trivial LWE encryption.

Then, the bootstrap operation can be defined as: $\text{CAdd}(B, \text{ExtractSample}(0, \text{BlindRotate}(L, R)))$.

```

def homomorphic_nand(b0: LWECiphertext, b1: LWECiphertext, constCipher: LWECiphertext, scale_lwe:
↪ LWECiphertext, txR: RLWECiphertext, bsk: BootstrapKey) -> LWECiphertext:
    """
    This function implements the nand operation between two inputs b0 and b1 homomorphically.
    The way it does it is by first evaluating the value of  $F(b0, b1) = \text{LWE}(-3) - b0 - b1$  and then
    ↪ running the bootstrap that implements the step function which evaluates the final nand
    ↪ function.
    """

```

*Remember that 0 is equivalent to F while 2 is equivalent to T here. The output will be an LWE
↪ ciphertext encrypting 0 or 2 as well.
constCipher is just the LWE Encryption of -3.
'''*

```
# calculating constCipher - b0 - b1
cc_minus_b0 = constCipher.sub_ciphertext(b0)
F_lwe = cc_minus_b0.sub_ciphertext(b1)
# call bootstrapping on this
nand_lwe = bootstrap(F_lwe, bsk, txR, scale_lwe)

return nand_lwe
```


Chapter 6

TFHE and TFHE-rs: Introduction

TFHE is a homomorphic encryption scheme. TFHE-rs is a rust library implemented by ZAMA that implements a variant of TFHE. TFHE follows the same methods and the ciphertexts mentioned in previous chapters. A few of its notable benefits are that it supports evaluation of small lookup tables using programmable bootstrapping, which it implements very efficiently and its runtime is fairly quick compared to other schemes. However, it has a few disadvantages. First, it can only work with small message sizes, meaning, the plaintext modulus (p) is fairly small). Second, the scale difference between p and q is huge (q is much larger), because of which, in vanilla TFHE, you have limited bit-widths to work with (usually 1 - 8 bits).

The TFHE-rs library attempts to and successfully solves this scale issue. On a higher level, does so by using multiple LWE ciphertexts to encrypt a single integer. It splits large integers (say a $u64$) into smaller blocks (say $u8$), and each LWE ciphertext encrypts one block. The sequence of these blocks is important to extract the final integer out, and to ensure that, an encoding is used to preserve that sequence. The library abstracts all the internal math of this sequence of LWE ciphertexts so that the user can focus directly on arithmetic.



Figure 6.1: LWE encryption of one block (image taken from the TFHE-rs Handbook).

The yellow blocks represent the a component (array of size n), and the grey block is the b component for LWE encryption




Figure 6.2: Encryption of larger scale integer using LWE blocks (each 'row' is an LWE block; image taken from the TFHE-rs Handbook)

6.1 TFHE: Setup and Introduction

TFHE, as stated above, is an encryption scheme which is also based on (G)LWE ciphertexts. Where it is different from other schemes is that TFHE implements a very efficient bootstrapping operation (orders of magnitude faster

than other schemes afaik), and this bootstrapping is also programmable, meaning, in addition to resetting the noise in the ciphertext, it can also evaluate a function on it and return the encrypted output of the function, which is why it is often referred to as *programmable bootstrapping (PBS)* or *functional bootstrapping*.

One thing to note is that multiplication is not natively available in TFHE (compared to something like CKKS, which allows multiplication by design). However, PBS can be used to perform multiplication, which converts it from an additive scheme to a fully homomorphic encryption scheme.

NOTE ABOUT NOTATION: In this and the coming chapters, a polynomial will be denoted by an uppercase letter, while a constant will be denoted by a lowercase letter. Vectors of elements will be denoted by bold letters based on their corresponding types (I will try my best to follow this bold thing, but I might forget, that will not be a problem because I will most definitely mention that that particular letter references a vector of a specific element).

6.1.1 Hard Lattice Problems: LWE and GLWE Assumptions

TFHE and most other homomorphic encryption schemes are based on a family of hard problems, the **LWE (Learning With Error)** problem. This is what ensures its security. Basically, it can be thought of as something similar to solving a system of noisy linear equations. This problem is of two types: a search problem and a decision problem.

LWE: Learning With Errors

Let:

- $n \in \mathbb{N}$, a positive integer called the LWE dimension,
- $q \in \mathbb{N}$, a positive integer, called the ciphertext modulus,
- $\mathbf{s} = (s_0, s_1, \dots, s_{n-1}) \in \mathbb{R}_q^n$ be a vector that is obtained from a distribution $\mathcal{D}$ known as the secret distribution. Conversely, $\mathbf{s}$ is called the secret,¹
- a distribution χ called the noise distribution.

Then, a sample from this LWE distribution that is defined using all these parameters is denoted by $LWE_{n,q,\chi,\mathcal{D}}$. Additionally, a point sampled from this distribution looks like:

$$(a, b = \sum_{i=0}^{n-1} a_i \cdot s_i + e_i) \in \mathbb{Z}_q^{n+1}$$

where, $\mathbf{a} = (a_0, a_1, \dots, a_{n-1}) \in \mathbb{Z}_q^n$ is a vector sampled from a uniform distribution $\mathcal{U}(\mathbb{Z}_q)$ ¹.

This is the LWE distribution, and as mentioned above, there can be two different problems made out of this:

- **Search Problem:** This problem is about trying to find out the secret vector $\mathbf{s}$ from a number of samples generated from $LWE_{n,q,\chi,\mathbb{D}}$.
- **Decision Problem:** In this problem, we try to distinguish between samples from the $LWE_{n,q,\chi,\mathbb{D}}$ and a uniform distribution $\mathcal{U}(\mathbb{Z}_q)$.

The security of encryption schemes that use either of these LWE problems as their basis depend of the values of q , n , distributions χ and $\mathcal{D}$. Additionally, the security of an LWE based ciphertext is denoted by something called as the *security factor* which is denoted by λ . For a scheme with a specific configuration to have a security level λ , breaking it should require atleast 2^λ operations. $\lambda = 128$ is considered to be a good value to ensure security of such systems.

¹Every element of the array is independently sampled from its distribution.

GLWE: General Learning With Errors

The problem with LWE is that it works with integers, and LWE ciphertexts also have a high expansion factor between plaintext and ciphertext (there is a large difference between magnitudes of p and q). There can be instances where one needs to work with more complex structures such as polynomials, and for that, the General Learning With Errors problem is more suitable, as instead of encrypting and working with integers, it works with polynomials in a ring $R_q[X]/(X^N + 1)$. Since it works with a polynomial, it uses more complex mathematical operations than integers. It is important to remember that in GLWE, the coefficients of polynomials are modulus q , and the polynomials themselves are modulus $(X^N + 1)$, meaning, their maximum degree is $N-1$. This means that after any arithmetic operation, specifically multiplication, the resulting polynomial is reduced to a polynomial of degree $N - 1$ by taking mod wrt $(X^N + 1)$, and the same goes for the coefficients where a modulus q operation is also applied on them (this is also done with LWE ciphertexts).

Let:

- $k \in \mathbb{Z}_+$, a positive integer, called the GLWE Dimension,
- $N \in \mathbb{Z}_+$, a positive integer, called the polynomial size,
- χ , the noise distribution,
- $\mathbb{D}$, the secret key distribution.
- $q \in \mathbb{Z}_+$, a positive integer, the ciphertext modulus

Let $\mathbf{S} = (S_0, S_1, \dots, S_{k-1})$ be the vector of polynomials called the secret, where the coefficient for each polynomial S_i is sampled from the secret distribution $\mathbb{D}$. Then, a sample from the GLWE distribution $\text{GLWE}_{q,k,N,\chi,\mathbb{D}}$ is defined as:

$$\text{GLWE}_{q,k,N,\chi,\mathbb{D}} = (\mathbf{A}, B = \sum_{i=0}^{k-1} A_i \cdot S_i + E) \in \mathcal{R}_{q,N}^{k+1}^2$$

where, $\mathbf{A} = (A_0, A_1, \dots, A_{k-1})$ is an array of k polynomials, where each polynomial's coefficient is sampled from a uniform distribution $\mathcal{U}(\mathbb{Z}_q)$ independently, and the noise polynomial E has its coefficients sampled independently from the noise distribution χ .

Like the LWE system, this also has two problems, which are the same as the LWE ones but it works on polynomials:

- **Search Problem:** Trying to find out the secret polynomial $\mathbf{S}$ from an arbitrary number of GLWE samples.
- **Decision Problem:** Trying to differentiate between samples from the RLWE distribution and points/polynomials sampled from the uniform $\mathcal{U}(\mathbb{R}_{q,N}^{k+1})$

Two important points:

- For GLWE, if $N = 1$, $k=n$, then it becomes equivalent to the LWE system.
- If $k = 1$, then the resulting GLWE system is called the RLWE system (it works with 1 polynomial instead of an array of polynomials).

6.2 FHE Ciphertext Types

6.2.1 LWE, RLWE & GLWE Ciphertexts

The LWE & RLWE ciphertexts are what we looked at in the previous chapters. The GLWE ciphertext is sort of a generalization of LWE. Think of GLWE (this isn't how it literally happens, but I think it's a good intuitive explanation) as sort of an RLWE encryption set but instead of the secret key and $\mathbf{A}$ being a single polynomial, they're an array of polynomials.

Given an encoded message polynomial $\tilde{M} \in R_q^n$, and a secret $\mathbf{S} = (S_0, S_1, \dots, S_{k-1}) \in R_{q,N}^k$, the GLWE encryption of the message is:

²Remember that bold uppercase letters represent an array of polynomials, while lowercase bold letters represent an array of scalars.

$$CT = ((A_0, A_1, \dots, A_{k-1}), B = \sum_{i=0}^{k-1} A_i \cdot S_i + \tilde{M} + E) = GLWE_{\mathbf{S}}(\tilde{M}) \in \mathcal{R}_{q,N}^{k+1}$$

The random elements are sampled the same way as mentioned above in the GLWE system.

Note that while ideally, the coefficients of the noise polynomial E should be sampled independently, it is not mandatory, but ideally, this property would be there. Additionally, the different coefficients can be sampled from different noise distributions $\mathcal{D}$.

6.2.2 Lev, RLev & GLev Ciphertexts:

These are a family of more complex ciphertexts that can be used for some applications downstream. They are actually a collection of GLWE ciphertexts, but they have two additional parameters, *a decomposition base* and *a decomposition level*. These ciphertexts do not encode a plaintext (encoded message), but rather, the cleartext raw message itself, and thus, assumes that the raw message is much smaller compared to the ciphertext modulus.

Given the GLWE secret key $\mathbf{S}$, the ciphertext modulus q , a decomposition level $l \in \mathbb{N}$ and a decomposition base $\mathcal{B} \in N$, the GLev ciphertext encrypting a cleartext message M (not encoded) is a collection of GLWE ciphertexts under the same configurations:

$$CT_{Glev} = GLew_{\mathbf{S}}^{\mathcal{B},l}(M) = (CT_0, CT_1, \dots, CT_{l-1}) \in \mathcal{R}_{q,N}^{l \cdot (k+1)}$$

where each ciphertext CT_i is an GLWE Ciphertext such that:

$$CT_i = GLWE_{\mathbf{S}}(\frac{q}{\mathcal{B}^{(i+1)}} \cdot M) \in \mathcal{R}_{q,N}^{k+1} \quad \forall \quad 0 \leq i < l$$

When $N = 1$ and $k = n$, the resulting ciphertext is called Lev (think of its a LWE equivalent for GLev), and when $N \neq 1$ and $k=1$, then its called an RLev (think of it as RLWE equivalent of GLev).

6.2.3 GSW, RGSW & GGSW Ciphertexts

On a high level, GGSW ciphertexts are a collection of GLev ciphertexts that use elements of secret key to define the GLev plaintexts. This sounds very vague but the mathematical definition itself does a good job of explaining what they are in my opinion. Given the decomposition base $\mathcal{B}$, the decomposition level l , and the GLWE secret key $\mathbf{S} = (S_0, S_1, \dots, S_{k-1})$, the GGSW ciphertext is defined as:

$$GGSW_{\mathbf{S}}^{\mathcal{B},l}(M) = (CT_0, CT_1, \dots, CT_k) \in \mathcal{R}_{q,N}^{(k+1)^2 \cdot l}$$

such that:

$$CT_i = GLew_{\mathbf{S}}^{\mathcal{B},l}(-S_i \cdot M) \in \mathcal{R}_{q,N}^{(k+1) \cdot l} \quad \forall \quad 0 \leq i \leq k$$

6.3 Encoding and Decoding

The encoding and decoding methods for TFHE are almost the same as that discussed above. However, there is a very small difference, that is, you can have padding bits in TFHE.

6.3.1 What does padding mean?

Generally, there are two things defined, p : plaintext modulus and q : ciphertext modulus, and to encode a message, you multiply it with $\Delta = \frac{q}{p}$. What happens in padding is that in addition to p and q , a value $\pi > 0$ called the padding size is defined. The value of Δ now changes to $\Delta = \frac{q}{p \cdot 2^\pi}$.



Figure 6.3: Representation of plaintext when padding is used. Here, padding is set to be 2 bits ($\pi = 1$).

Image taken from the TFHE-rs Handbook

6.3.2 Why is padding used?

While it might look like padding wastes valuable space, it is important for applications such as PBS. This is because padding allows us to implement PBS based table lookups for non-negacyclic functions. Without it, you'd be stuck to working only with negacyclic functions, while in reality, many functions are not negacyclic. Having a padding bit, and most importantly, having the padding bit known (as in what value it holds) allows one to go for non-negacyclic functions. Typically, the value in the padding bit is kept to zero.

6.3.3 Carry Bits

We want the value in the padding bit(s) to be known (generally 0). But during some arithmetic operations, it is possible that the resulting message has a higher bit width and it goes beyond the plaintext modulus p and modifies the value of the padding bit. For example, assume that $p = 2^2$, and you have two messages, $m_1 = 2 = 01_2$, and $m_2 = 3 = 11_2$. Adding them both gives $5 = 101_2$, which would definitely change the value of the padding bit from 0 to 1, and that would result in an inaccurate PBS operation. To prevent this unwanted (unless its explicitly done) updating, carry bits are used. Two things are defined:

- Message Modulus, **msg**
- Carry Modulus, **carry**

such that $\mathbf{carry} \cdot \mathbf{msg} = p$. The message modulus sets the range for the actual cleartext message.

This ensures that even if the result of the output goes beyond the message modulus, it can never edit the padding bit, ensuring accurate bootstrapping. Once the carry space becomes full, a PBS operation be done, to reduce the message modulus wrt **msg**, which resets the carry bits.

(Thoughts to myself: Okay if you eventually have to take modulus once the carry bit is full, can you not also do the same for when the padding space gets edited? That removes the need for shrinking the actual message space. I can think of the following reasons:)

- If the parameters are set correctly, then this situation never arises, unless something crazy happens. This is more of a foolproof safety check.
- This PBS operation is hard and you want to delay it as long as possible.

6.4 Encoding Booleans

Booleans (1/0; T/F) can be represented using integers, 0 for False/0, 2 for True/1 for example, and then they can be encoded the same way as ints.

6.5 Simple Operations on TFHE Ciphertexts

Simple arithmetic operations such as addition and multiplication with a scalar and addition with another ciphertext is done the same way as stated in the previous chapters for GLWE ciphertexts. For GGWE & GLev ciphertexts, since they're essentially sequences of GLWE ciphertexts, these operations from GLWE can be trivially done by performing those operations on the corresponding GLWE building blocks.

6.6 Key Switching

For any operation that needs to be done between two ciphertexts, it is important that they're encrypted using the same secret key. Ciphertexts encrypted under different secret keys can not be correctly added, multiplied etc even if their shapes are the same. To be able to do things like these, the key switching operation is done. If you have a ciphertext $CT_{in} = \text{GLWE}_{\mathbf{s}_{in}}(M)$, using key switching, you can get another ciphertext $CT_{out} = \text{GLWE}_{\mathbf{s}_{out}}(M)$, that is encrypted under a different key compared to the input ciphertext. Note that this 'switched' key can have different parameters such as N , q etc than the input secret key. There are many different ways/algorithms for key switching, and all of

them requires something called as the keyswitching keys. However, intuitively, keyswitching keys are encryptions of the input ciphertext's key $\mathbf{S}_{\text{in}}$ using the output secret key $\mathbf{S}_{\text{out}}$. Note that the previous chapters talk about switching keys between different ciphertexts. This section will take about other algorithms, specific to TFHE.

6.6.1 Radix Decomposition

Radix Decomposition is a part of most key switching operations. Radix decomposition is defined for an integer $x \in \mathbb{Z}_q$. For a given decomposition base $\mathcal{B}$ and decomposition level l , such that $\mathcal{B}^l$ divides q . The idea is to decompose x into a vector $(x_1, x_2, \dots, x_l)$, such that:

$$\langle (x_1, x_2, \dots, x_l), (\frac{q}{\mathcal{B}}, \frac{q}{\mathcal{B}^2}, \dots, \frac{q}{\mathcal{B}^l}) \rangle = \sum_{i=1}^l x_i \cdot \frac{q}{\mathcal{B}^i} = \lfloor x \cdot \frac{q}{\mathcal{B}^l} \rfloor \cdot \frac{q}{\mathcal{B}^l} \in \mathbb{Z}_q$$

Note that $\langle \cdot, \cdot \rangle$ represents the dot product between the two. (Also, while the above equality might look daunting, the right hand side of the equality should equate to x or something extremely close to x if everything is correct and the parameters are right.)

It is implemented in TFHE-rs using an algorithm called $\text{BalDecomp}^{\mathcal{B},l}(x)$

6.6.2 LWE and GLWE Key Switch

LWE Key Switch

The LWE Key Switch operation takes as input an LWE ciphertext encrypting a plaintext under the a secret key $\mathbf{s} = (s_0, s_1, \dots, s_{n-1})$, and outputs a GLWE ciphertext encrypting the same plaintext under a secret key $\mathbf{S} = (S_0, S_1, \dots, S_{k-1})$. In the case when $N = 1$, then the output of this key switch operation becomes an LWE ciphertext. Additionally, this function also requires a keyswitch key, which is an array of GLew ciphertexts where each element encrypts a single element of $\mathbf{s}$:

$$\text{KSK} = \{CT_i : \text{GLew}_{\mathbf{S}}^{\mathcal{B},l}(s_i)\}_{0 \leq i \leq n-1}$$

Algorithm to perform this keyswitch is as follows:

Algorithm 10: $\text{CT} \leftarrow \text{KS}(\text{ct}, \text{KSK})$	
Context:	$\begin{cases} \mathbf{s} = (s_0, \dots, s_{n-1}) : \text{the input LWE secret key} \\ \mathbf{S} = (S_0, \dots, S_{k-1}) : \text{the output GLWE secret key} \end{cases}$
Input:	$\begin{cases} \text{ct} = (a_0, \dots, a_{n-1}, b) \in \text{LWE}_{\mathbf{s}}(\tilde{m}) \\ \text{KSK} = \{\overline{\text{CT}}_i \in \text{GLew}_{\mathbf{S}}^{\mathcal{B},l}(s_i)\}_{0 \leq i \leq n-1} : \text{keyswitching key from } \mathbf{s} \text{ to } \mathbf{S} \end{cases}$
Output:	$\text{CT} \in \text{GLWE}_{\mathbf{S}}(\tilde{m})$
1	$\text{CT} \leftarrow (0, \dots, 0, b) - \sum_{i=0}^{n-1} \langle \overline{\text{CT}}_i, \text{BalDecomp}^{\mathcal{B},l}(a_i) \rangle$
2	return CT

Figure 6.4: LWE Keyswitching algorithm, image taken from the TFHE-rs handbook.

GLWE Key Switch

The GLWE keyswitch takes as input a GLWE ciphertext encrypted under a GLWE key $\mathbf{S}_{\text{in}}$ and outputs another GLWE ciphertext encrypting the same message under a different GLWE key $\mathbf{S}_{\text{out}}$. The keyswitch key here is a sequence of GLew ciphertexts encrypting individual polynomials of $\mathbf{S}_{\text{out}}$ using the $\mathbf{S}_{\text{in}}$.

$$\text{KSK} = \{\text{CT}_i : \text{GLew}_{\mathbf{S}_{\text{in}}}^{\mathcal{B},l}(S_{\text{out},i})\}_{0 \leq i \leq k_{\text{in}}-1}$$

The algorithm for this keyswitch operation is as follows:

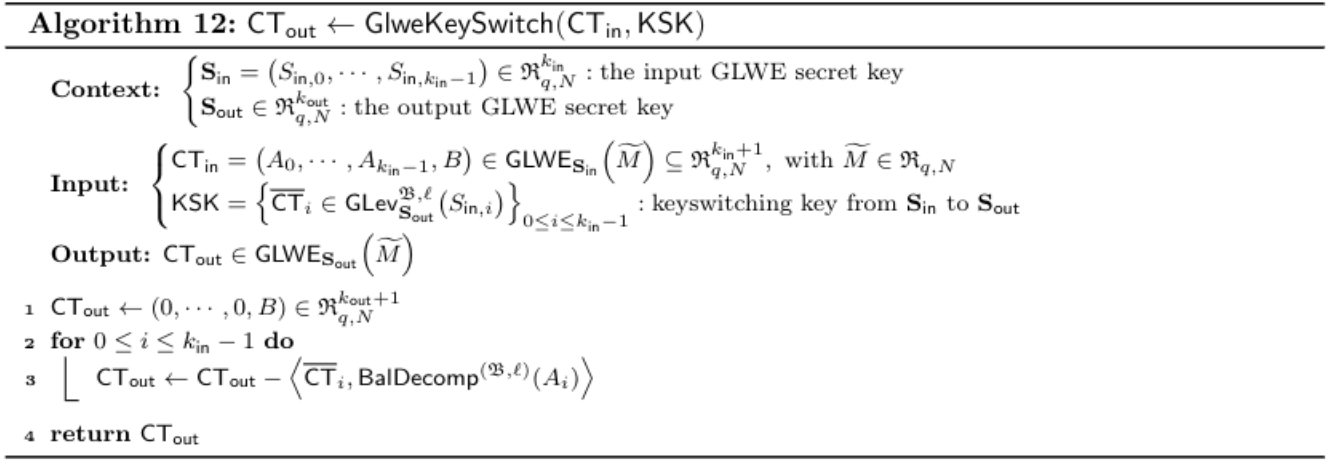


Figure 6.5: GLWE keyswitching algorithm, taken from the TFHE-rs handbook.

6.6.3 Packing Key Switch

Packing is the process of taking many LWE ciphertexts and 'packing' them into one GLWE ciphertext, this works because a GLWE ciphertext encrypts a polynomial. Basically, if you have LWE ciphertexts that encrypt a collection of numbers, say $(l_0, l_1, \dots, l_{N-1})$, then the packing will result in a GLWE ciphertext that encrypts a polynomial $l_0 + l_1x + \dots + l_{N-1}x^{N-1}$.

The packing key switch operation takes a collection of $\alpha < N$ LWE ciphertexts $\{ct_0, ct_1, \dots, ct_{\alpha-1}\}$, encrypted under an LWE key $\mathbf{s}$ and a set of α indices for these ciphertexts, $\alpha = \{i_j\}_{j=0}^{\alpha-1}$ and returns a GLWE ciphertext that encrypts a polynomial where the coefficients are the messages encrypted in the input LWE ciphertexts.

Here, the keyswitch key again is the GLWE encryption of the LWE key $\mathbf{s}$ using the GLWE key $\mathbf{S}$. The algorithm to perform this operation is given below:




Figure 6.6: Algorithm to perform packing key switching. Image taken from the TFHE-rs handbook.

6.7 Programmable Bootstrapping (PBS): Building Blocks

6.7.1 Modulus Switching

The modulus switch operation switches the ciphertext modulus of an encrypted ciphertext from existing q to a new modulus w . If you have a ciphertext that encrypts an encoded message $\tilde{m}$ (encoded using q), the output of this

operation returns a ciphertext that is encoded in w : $\lfloor \tilde{m} \cdot \frac{w}{q} \rfloor$. For TFHE's PBS, w is set to be $2N$. There are 3 types of modulus switching operations, which are controlled by the parameter `pfailMitigationTechnique`:

- `pfailMitigationTechnique = 0`: No mitigation technique used, no drift mitigation key needed, classical modulus switching is used.
- `pfailMitigationTechnique = 1`: Mean compensation technique used for which no drift mitigation key is needed, after that, classical modulus switching is used.
- `pfailMitigationTechnique = DMK`: Drift mitigation technique is used before performing modulus switching. Drift Mitigation Key (DMK) needed.

The modulus switching algorithm:




Figure 6.7: Modulus Switching Algorithm, image taken from the TFHE-rs handbook.

Mean Compensation

Decrypting a ciphertext produces a noisy version of the encoded ciphertext (because it returns $m + e$, and not just m). This is then decoded to get the original message. However, there is a chance that the decryption and decoding fail, which is called the failure probability. This probability is directly influenced by the magnitude of noise in the ciphertext (higher the noise, higher this probability). Also, during the process of modulus switching, the noise comes from rounding and the resulting rounding errors. The drift vector is just a vector of these rounding errors that happen during modulus switching.

The mean compensation technique aims at reducing this noise by 'adjusting' the mean of this drift vector to zero.

Algorithm 15: $\text{ct}_{\text{out}} \leftarrow \text{MeanCompensation}(\text{ct}_{\text{in}})$

Context: $\mathbf{s} = (s_0, \dots, s_{n-1}) \in \{0, 1\}^n$: the input binary LWE secret key

Input: $\text{ct}_{\text{in}} \in \text{LWE}_{\mathbf{s}}(\tilde{m}) \subseteq \mathbb{Z}_q^{n+1}$, where $\tilde{m} = \Delta \cdot m$

Output: $\text{ct}_{\text{out}} \in \text{LWE}_{\mathbf{s}}(\tilde{m})$

```
1  $\text{ct}_{\text{in}} = (a_{\text{in},0}, \dots, a_{\text{in},n-1}, b_{\text{in}})$ 
2  $\tilde{\mu} \leftarrow -\sum_{j=0}^{n-1} \left( \left\lfloor a_{\text{in},j} \frac{2N}{q} \right\rfloor \frac{q}{2N} - a_{\text{in},j} \right) // 2$ 
3  $\text{ct}_{\text{out}} \leftarrow (a_{\text{in},0}, \dots, a_{\text{in},n-1}, b_{\text{in}} - \tilde{\mu})$ 
4 return  $\text{ct}_{\text{out}}$ 
```

Figure 6.8: Mean Compensation Algorithm, image taken from the TFHE-rs handbook.

Drift Mitigation

Drift Mitigation is an alternative to Mean Compensation for reducing noise during the modulus switching operation. The idea here is that, the rounding error is dependent on the input mask ($\mathbf{a}$). Adding a 0 LWE ciphertext to the input randomizes this mask and thus the rounding error. While this also adds a little extra noise (as you're adding two LWE ciphertexts), this is okay as the rounding noise reduction is much higher than the added noise. To perform this operation, an extra key, called the Drift Mitigation Key (DMK) is required, which are essentially fresh LWE encryptions of 0. $\text{DMK} = \{\text{ct}_z = \text{LWE}(0)\}_{z \in [1, Z]}$.

Drift Mitigation Algorithm:

Algorithm 16: $\text{ct}_{\text{out}} \leftarrow \text{DriftMitigation}(\text{ct}_{\text{in}}, \text{DMK})$

Context: $\begin{cases} \mathbf{s} = (s_0, \dots, s_{n-1}) \in \mathbb{Z}^n : \text{the input binary LWE secret key} \\ T, r, Z : \text{the drift parameters} \end{cases}$

Input: $\begin{cases} \text{ct}_{\text{in}} \in \text{LWE}_{\mathbf{s}}(\tilde{m}) \subseteq \mathbb{Z}_q^{n+1}, \text{ where } \tilde{m} = \Delta \cdot m \\ \text{DMK the drift mitigation key} \end{cases}$

Output: $\text{ct}_{\text{out}} \in \text{LWE}_{\mathbf{s}}(\tilde{m})$

```
1 index = 0
2  $\text{ct}_{\text{in}} = (a_{\text{in},0}, \dots, a_{\text{in},n-1}, b_{\text{in}})$ 
3  $\tilde{\mu} \leftarrow \left( \left\lfloor b_{\text{in}} \frac{2N}{q} \right\rfloor \frac{q}{2N} - b_{\text{in}} \right) - \sum_{j=0}^{n-1} \left( \left\lfloor a_{\text{in},j} \frac{2N}{q} \right\rfloor \frac{q}{2N} - a_{\text{in},j} \right) \cdot \frac{1}{2}$ 
4  $\tilde{\sigma}^2 \leftarrow \sum_{j=0}^{n-1} \left( \left\lfloor a_{\text{in},j} \frac{2N}{q} \right\rfloor \frac{q}{2N} - a_{\text{in},j} \right)^2 \cdot \frac{1}{4}$ 
5  $T_{\text{best}} = |\tilde{\mu}| + r \cdot \tilde{\sigma}$ 
6 index = 0
7 while  $|\tilde{\mu}| + r \cdot \tilde{\sigma} \geq T$  and index < Z do
8   index  $\leftarrow$  index + 1
9    $\text{ct}_{\text{DriftMitigation}} = (a_0, \dots, a_{n-1}, b) \leftarrow \text{Add}(\text{ct}_{\text{in}}, \text{ct}_{\text{index}})$ 
10   $\tilde{\mu} \leftarrow \left( \left\lfloor b \frac{2N}{q} \right\rfloor \frac{q}{2N} - b \right) - \sum_{j=0}^{n-1} \left( \left\lfloor a_j \frac{2N}{q} \right\rfloor \frac{q}{2N} - a_j \right) \cdot \frac{1}{2}$ 
11   $\tilde{\sigma}^2 \leftarrow \sum_{j=0}^{n-1} \left( \left\lfloor a_j \frac{2N}{q} \right\rfloor \frac{q}{2N} - a_j \right)^2 \cdot \frac{1}{4}$ 
12   $T_{\text{index}} = |\tilde{\mu}| + r \cdot \tilde{\sigma}$ 
13  if  $T_{\text{index}} < T_{\text{best}}$  then
14     $T_{\text{best}} = T_{\text{index}}$ 
15    indexbest = index
16 if index = Z and  $T_{\text{index}} > T$  and indexbest  $\neq$  0 then
17    $\text{ct}_{\text{DriftMitigation}} \leftarrow \text{Add}(\text{ct}_{\text{in}}, \text{ct}_{\text{index}_{\text{best}}})$ 
18 return  $\text{ct}_{\text{out}} = \text{ct}_{\text{DriftMitigation}}$ 
```

Figure 6.9: Drift Mitigation Algorithm, image taken from the TFHE-rs handbook.

6.7.2 External Product

External product essentially is the ciphertext-ciphertext multiplication. Like shown in the previous chapters, where we performed multiplication between a GSW and an RLWE ciphertexts, here, we generalize it to GGSW and GLWE ciphertexts. Instead of using the base-p representation for RLWE ciphertext, we use the BalDecomp (Balanced Radix decomposition).




Figure 6.10: External Multiplication Algorithm, image taken from the TFHE-rs handbook.

6.7.3 CMux

The CMux operation is also extremely similar to what we saw previously, except that now, it works with more general GLWE and GGSW ciphertexts.




Figure 6.11: CMux Algorithm for TFHE, image taken from the TFHE-rs handbook.

6.7.4 Blind Rotate

The blind rotate operation, just like the other operations, is similar in how its done in the previous chapters, except that the ciphertexts are the general versions. The bootstrapping key is a sequence of GGSW ciphertexts that encrypts each component of the LWE encryption key.

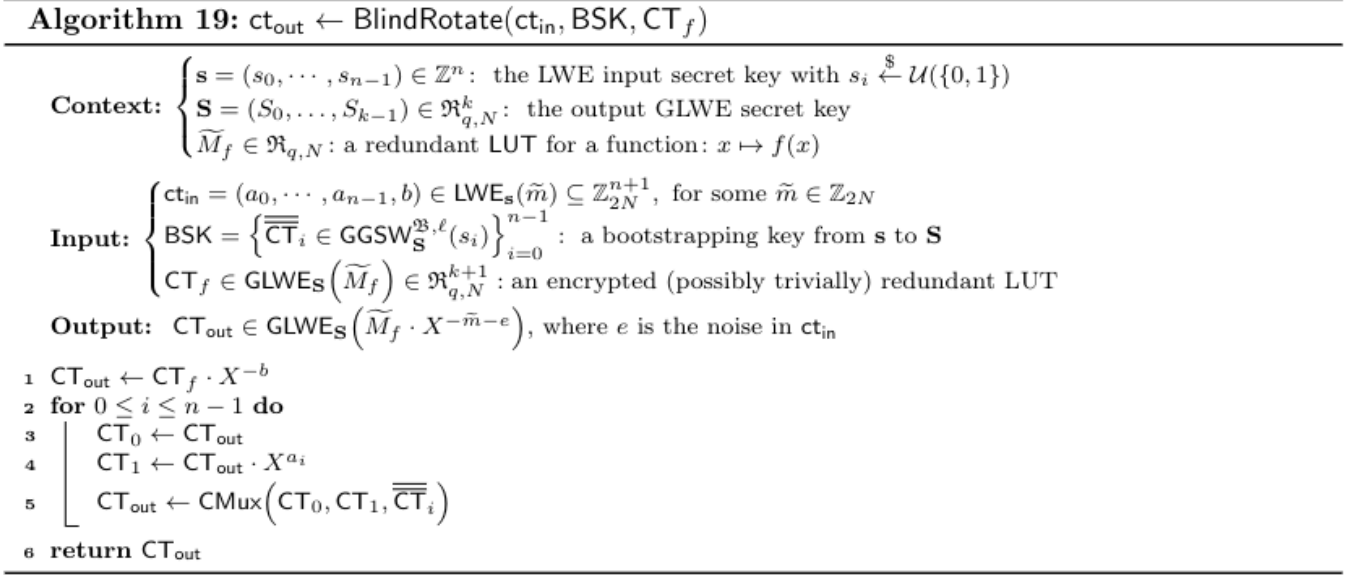


Figure 6.12: BlindRotate Algorithm, image taken from the TFHE-rs handbook.

6.8 Programmable Bootstrapping

The programmable bootstrapping in TFHE achieves two things:

- It refreshes noise of the output ciphertext (noise of the output is independent of the input ciphertext).
- It allows one to evaluate a function on the input ciphertext and return the output of that function encrypted as an LWE ciphertext.

While the PBS implementation is similar to the ones in the previous chapters, one big difference is its ability to evaluate a function, for which, it takes as input, something known as a redundant LUT.

Algorithm 20: $\text{ct}_{\text{out}} \leftarrow \text{PBS}(\text{CT}_f, \text{ct}_{\text{in}}, \text{BSK}, \text{pfailMitigationTechnique})$

Context: $\begin{cases} \mathbf{s} = (s_0, \dots, s_{n-1}) \in \mathbb{Z}^n : \text{the input binary LWE secret key} \\ \mathbf{S}' = (S'_0, \dots, S'_{k-1}) \in \mathfrak{R}_{q,N}^k : \text{a GLWE secret key where } S'_i = \sum_{j=0}^{N-1} s'_{i \cdot N+j} X^j \text{ for all } i \\ \mathbf{s}' = (s'_0, \dots, s'_{kN-1}) \in \mathbb{Z}_q^{kN} : \text{the output LWE secret key} \\ \widetilde{M}_f \in \mathfrak{R}_{q,N} : \text{a redundant LUT for a function } f \\ \ell : \text{the decomposition level} \\ \mathfrak{B} : \text{the decomposition base} \end{cases}$

Input: $\begin{cases} \text{CT}_f \in \text{GLWE}_{\mathbf{S}'}(\widetilde{M}_f) \subseteq \mathfrak{R}_{q,N}^{k+1} \\ \text{ct}_{\text{in}} \in \text{LWE}_{\mathbf{s}}(\widetilde{m}) \subseteq \mathbb{Z}_q^{n+1}, \text{ where } \widetilde{m} = \Delta \cdot m \\ \text{BSK} : \text{a bootstrapping key from } \mathbf{s} \text{ to } \mathbf{S}' \\ \text{pfailMitigationTechnique} \in \{0, 1, \text{DMK}\} \end{cases}$

Output: $\text{ct}_{\text{out}} \in \text{LWE}_{\mathbf{s}'}(\widetilde{f(m)})$ if we respect the requirements of Theorem 12

1 $\text{ct}_{\text{MS}} \leftarrow \text{ModulusSwitch}(\text{ct}_{\text{in}}, 2N, \text{pfailMitigationTechnique})$
2 $\text{CT} \leftarrow \text{BlindRotate}(\text{ct}_{\text{MS}}, \text{BSK}, \text{CT}_f)$
3 $\text{ct}_{\text{out}} \leftarrow \text{SE}(\text{CT}, 0)$
4 **return** ct_{out}

Figure 6.13: PBS Algorithm, image taken from the TFHE-rs handbook.

Let $\mathbf{s} = (s_0, s_1, \dots, s_{n-1})$ be the LWE encryption key for the input, and $\mathbf{S}' = (S'_0, S'_1, \dots, S'_{k-1})$ be the GLWE encryption key, and the corresponding flattened LWE secret key $\mathbf{s}' = (s'_0, s'_1, \dots, s'_{kN-1})$ s.t $S'_i = \sum_{j=0}^{N-1} s'_{i \cdot N+j} \cdot X^j$. The BSK is the associated bootstrapping key from $\mathbf{s}$ to $\mathbf{S}'$. For an LWE ciphertext encrypting an encoded message $\widetilde{m} = m \cdot \Delta$ for some $m \in \mathbb{Z}_{2p}$ for some p dividing N and $\Delta = \frac{q}{2p}$. $\widetilde{M}_f$ is the $\frac{N}{p}$ -redundant LUT for a function $f : \mathbb{Z}_p \rightarrow \mathbb{Z}_q$ and let CT_f be a trivial RLWE encryption of $\widetilde{M}_f$. These are the inputs to the bootstrap operation alongside the LWE ciphertext.

Redundant Lookup Table (LUT)

A redundant LUT is an LUT that encodes a function f (encodes, not encrypts), and the entries (outputs I think) of this function are stored as coefficients of a polynomial in $\mathfrak{R}_{q,N}$. The redundancy comes from storing entries $f(i)$ r times with a certain shift. The term mega-cell is used to denote each block of redundant values in an r -redundant lookup table.

$$\widetilde{M}_f = X^{-r/2} \sum_{i=0}^{N/r-1} X^{i \cdot r} \cdot \left(\sum_{j=0}^{r-1} \widetilde{f(i)} \cdot X^j \right)$$

$\widetilde{f(i)}$ is an encoding of the function output $f(i)$. Another way to write this LUT is just using the mega cell values:

$$\left[\widetilde{f(0)}, \widetilde{f(1)}, \dots, \widetilde{f\left(\frac{N}{r}-1\right)} \right]$$

Note that this is the basic PBS operation. There are other types of PBS operations that have their own characteristics.

Chapter 7

PBS Implementation In Depth

Please note that this chapter might contain some redundancy as I discuss the points already discussed in the previous chapters, however, here I try to go into a little more detail and implement some more PBS operations. Programmable Bootstrapping, as studied above, does two things:

- Reduces/refreshes noise in a polynomial.
- Evaluates a function (in TFHE).

However, there are things that should be taken care of when PBS is being implemented for a specific function. The most important thing that needs to be remembered is that it takes as input an LWE ciphertext and outputs an LWE ciphertext too, and is traditionally, a univariate function. For bootstrapping, you need something called as the **Redundant Lookup Table** is very important.

7.1 Redundant Lookup Table/TLU

Lets take an example of a bootstrapping operation that tries to remove noise during the decryption process. Meaning, the input to the function is $\Delta \cdot m + e \in \mathbb{Z}_q$, and it outputs $m \in \mathbb{Z}_p$. However, $p < q$, and thus, a number of values of $\Delta \cdot m + e$ will correspond to the same m . This means that there is a redundancy.




Figure 7.1: Image showing transformation and redundancy. Taken from TFHE-Deep Dive, part 4

The collection of cell blocks that correspond to the same value are called mega-cases. Here, one thing to notice is that the case for 0 is split, with half of it being on either ends. This is because at $m = 0$, the noise can put it on either side of the number line, positive noise moves it left while negative noise moves it right.¹

¹Assuming that its modeled as unsigned integer and the first half is used for positive while the second half is used for negative integers, as is the case in TFHE-rs.

7.1.1 Now why is this important?

As studied earlier, bootstrapping is evaluated using rotation (blind-rotation) of a polynomial to be exact. After which, you extract value from a specific place, which is pre-defined, for example, in the previous chapters, the coefficient of x^0 was extracted which was the output of the bootstrapping operation. The TLU defined above is actually modelled as a (trivially encrypted) polynomial with the coefficients being the elements of the output of the TLU set up such that rotating this LUT ciphertext brings the correct output from one of the many redundant positions in the mega case to the position that is being observed, which gives us the correct answer.

7.2 Negacyclic Property and PBS

Here, we talk about the usable domains and how does that get affected by the negacyclic property of the ring. Remember that, a negacyclic ring $\mathbb{Z}_q/(X^N + 1)$ has an order of $2N$, meaning, $X^N \equiv -1$ and $X^{2N} \equiv 1$. And therefore, for any polynomial from this ring, you need to rotate it, by $2N$ to get the original polynomial back. Also, rotating a polynomial $f(x)$ by N , gives you its negation $-f(x)$. Therefore, the actual domain of the rotation input becomes $2N$ and not N . You dont want to push further because after that it just loops around. However, the latter half of the rotation outputs is just the negation of the former half (think: $\text{rotate}(i, F) = -\text{rotate}(i+N, f)$), and you cant change that.

The LUT polynomial also belongs to the ring and thus, also follows the same properties. Now if this was used to evaluate a function that is negacyclic, meaning $f(x+N) = -f(x)$, then the two things 'fit' well with each other, and you can use the entire domain of $2N$, but if it isnt, then it becomes an issue, because in the rotation, the only thing you can really control is the first half, upto N . And since you need to do a modulus switch from q to $2N$, ensuring that post switch, the input stays within N ($[-N/2, N/2]$), you can only use half of your plaintext domain.

7.3 Sign Operation Using PBS

This section describes implementing the sign operation using bootstrapping. The sign function is given as:

$$\text{sign}(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{otherwise} \end{cases} \quad (7.1)$$

The most important thing for any PBS operation that evaluates any function is to correctly create/set up the TLU polynomial. Once the polynomial is set up correctly, the actual bootstrapping portion remains the same, first, homomorphically rotating the polynomial based on the input, followed by a sample extract extracting the coefficient corresponding to x^0 .

T: The TLU polynomial

Consider a polynomial:

$$\mathbf{T} = 0 + 0x + 0x^2 + \dots + 0x^{N/2-1} - x^{N/2} - x^{N/2+1} - \dots - x^{N-1}$$

Try rotating this polynomial:

$$\begin{aligned} \text{Rotate}(1, T) &= 1 + 0x + \dots + 0x^{N/2} - x^{N/2-1} - \dots - x^{N-1} \\ \text{Rotate}(N/2, T) &= 1 + x + \dots + x^{N/2-1} - 0x^{N/2} - \dots - 0x^{N-1} \\ \text{Rotate}(N, T) &= 0 + 0x + \dots + 0x^{N/2-1} + x^{N/2} + \dots + 1x^{N-1} \\ \text{Rotate}(-1, T) &= 0 + 0x + 0x^2 + \dots + 0x^{N/2-2} - x^{N/2-1} - x^{N/2} - \dots - x^{N-2} - 0x^{N-1} \end{aligned}$$

Now for each of the rotate results, if you extract the first coefficient:

$$\begin{aligned} \text{ExtractSample}(0, \text{Rotate}(1, T)) &= 1 \\ \text{ExtractSample}(0, \text{Rotate}(N/2, T)) &= 1 \\ \text{ExtractSample}(0, \text{Rotate}(N, T)) &= 0 \\ \text{ExtractSample}(0, \text{Rotate}(-1, T)) &= 0 \\ \text{ExtractSample}(0, \text{Rotate}(0, T)) &= 0 \end{aligned}$$

This can be generalized as:

$$\text{ExtractSample}(0, \text{Rotate}(i, T)) = \begin{cases} 1 & \text{if } 0 < i \leq N/2 \\ 0 & \text{if } -N/2 < i \leq 0 \end{cases} \quad (7.2)$$

This polynomial can be used as TLU for our range of polynomials. This will work for our range. As we know, this is not a negacyclic polynomial, so, we can only use $\frac{p}{2}$ values, which means, the input will be in the range $[\frac{-p}{4}, \frac{p}{4})$. On doing modulus scaling, this becomes:

$$[\frac{-p}{4}, \frac{p}{4}) = [\frac{-p}{4} \cdot \Delta, \frac{p}{4} \cdot \Delta) = [\frac{-N}{2}, \frac{N}{2})$$

In this range, the sign will be correctly defined using this polynomial.

A few examples ($N = 1024$, $p = 16$, $q = 2^{32}$):

- Consider you want to find the sign of $x = 3$. On encoding and scheme switching, you get:

$$\frac{q}{p} \cdot 3 \cdot \frac{2N}{q} = \frac{3N}{8}$$

Which is smaller than $N/2$ and hence, on performing rotation with this input and extracting the coefficient, we get 1.

- Consider $x = -2$. On encoding and switching, you get:

$$\frac{q}{p} \cdot -2 \cdot \frac{2N}{q} = \frac{-N}{4}$$

Which is larger than $-N/2$, and hence, on performing rotation with this input and extracting the coefficient, we get 0.

```
# T generation
# generating the polynomial
tx_coeffs = np.zeros(config.N, dtype=np.int32)
tx_coeffs[config.N//2:] = -1
tx = Polynomial(N=config.N, coeffs=tx_coeffs)
# encoding tx
encode_multiplier = encode_plaintext(1, lweconfig)
# generating a trivial RLWE polynomial for tx
tx_rlwe = create_trivial_rlwe_ciphertext(tx.polynomial_const_multiply(encode_multiplier.m), config)

# bootstrapping function
def sign_bootstrap(i: LWECiphertext, bsk: BootstrapKey, tx_cipher: RLWECiphertext):
    """
    This function performs PBS which evaluates the sign function. This will return an LWE ciphertext
    ↪ encrypting 0 if the input was negative or 0, 1 otherwise.
    """
    # need to implement coeff(0, blindrotate(i, tx_cipher))
    # blind rotate tx_cipher
    rotated_tx = bsk.blindrotate(tx_cipher, i)
    # extract the 0th coefficient
    coeff_lwe = extract_sample(0, rotated_tx)
    return coeff_lwe
```

7.4 Bivariate PBS: Max using PBS

While PBS traditionally only takes a single LWE ciphertext as input, there might be situations where you want to evaluate a PBS on two LWE ciphertexts, for example, calculating min, max etc, where it might not be trivial to reduce them into a single ciphertext. In these situations, you can use a *packing trick*² to implement this.

What is this packing trick?

We want to somehow convert the two LWE ciphertexts into a single LWE ciphertext, such that both the ciphertexts are intact in this new *packed* polynomial. One way to do this is the following:

- You have the regular parameters p , q , and n for LWE.
- You create another parameter, lets called it p_b , such that $p_b \mid 2p$. It should be small enough such that atleast two p_b s can fit in p . This p_b is what actually defines the range of the cleartext now.
- Encode and encrypt the cleartexts using the actual plaintext modulus.
- Pack them using the following relationship. For two LWE ciphertexts, a and b :

$$p = a \cdot p_b + b$$

- For creating the LUT³, unpack the ciphertexts, by doing

$$\begin{aligned} a &= p/p_b \\ b &= p \% p_b \end{aligned}$$

This allows for bivariate PBS. A few things to note with this (these might be a little specific to TFHE-rs's core-crypto):

- Since core-crypto works with unsigned integers, the way signed integers are represented is that the first half is used for positive signed values and the second half of the unsigned range is used to represent corresponding negative signed values. And since here we are creating our values mod p_b , before converting the signed cleartext to unsigned, you should take a modulus wrt p_b .

```
// RUST CODE TFHE-rs
// --- ENCODING MESSAGE ---
let p = 1u64 << 8; // plaintext modulus = 2^8
let p_1 = 1u64 << 4; // this will actually define the available range for the clear messages. The idea
↳ is that the 2 LWE ciphertexts will be packed into one and then in the PBS, they will be extracted to
↳ allow for bivariate PBS. This allows for operations such as max, min etc.
let delta = (1u64 << 63) / p;

let msg1_clear = 7i8; // Im using 8 bit right now because with this noise level it works well. I will
↳ move up if needed but for now, im using 8 bit as it anyway is working for our model.
let msg2_clear = -3i8;
let msg1_clear_reduced = (msg1_clear as i64).rem_euclid(p_1 as i64) as u64;
let msg2_clear_reduced = (msg2_clear as i64).rem_euclid(p_1 as i64) as u64;

// MAX PBS ACCUMULATOR
let max_accumulator: GlweCiphertextOwned<u64> = generate_programmable_bootstrap_glwe_lut(
    polynomial_size,
    glwe_dimension.to_glwe_size(),
```

²Please be aware that I got this packing trick from claude, which mentions that this is what is being used in bivariate PBS in TFHE-rs (bivariate PBS functions exist in higher level APIs), but I haven't checked the source code to confirm this, but the method seems to be working.

³In the previous section, we created the LUT polynomial ourselves, but this one is non-trivial, and thus, we will use TFHE-rs's built in accumulator, which takes logic in plaintext and then creates the LUT polynomial dynamically. This is only used for creating the LUT and not when working with ciphertexts, and thus, it can use operations such as divide etc.


```

p as usize,
ciphertext_modulus,
delta,
|x: u64| { // this is 128 because p is  $2^8 = 256$ , and the latter half represents negative values
    // unpack
    let a = x / (p_1 as u64);
    let b = x % (p_1 as u64);
    let sa: i64 = if a < p_1 / 2 { a as i64 } else { a as i64 - p_1 as i64 }; // taking mod
    ↪ centered around 0 wrt p_1 to ensure that we get the correct values as they were
    ↪ encoded that way.
    let sb: i64 = if b < p_1 / 2 { b as i64 } else { b as i64 - p_1 as i64 };

    let m = sa.max(sb);
    m.rem_euclid(p as i64) as u64
},
);

```

Chapter 8

Max from a TFHE encrypted array

In this chapter, I'm exploring ways to find max from an array of LWE ciphertexts, to use for the TRF model project. I will be trying to find ways to find argmax/max out of an array. This will include implementing relevant research papers, using sorting algorithms etc, and comparing their speeds.

8.1 A non-comparison oblivious sort and its application to private k-NN

This paper implements a sorting algorithm for sorting a list of LWE ciphertexts that are 'packed' into an RLWE ciphertext by an operation called as **Public Functional Key Switching (PFKS)**. Additionally, it also extends it by using the developed sorting algorithm to applications in k Nearest Neighbors by finding the k smallest elements from a list of LWE ciphertexts. For the sorting algorithm, the paper uses something called as **Counting sort**, which is a sorting algorithm that is not based on the methods that use comparisons. The benefit of this, alongside it not using comparisons is that instead of the worst case $O(n \log n)$ complexity, it has a worst case complexity of $O(n + k)$ (n: input size; k: range of values). The authors of this paper have provided the code for the algorithm they propose in this github repository.

Algorithm 1: Counting Sort

Input : An array $[m_0, \dots, m_{p-1}]$ of length p
Output: The sorted array

```
1  $C \leftarrow [0, \dots, 0]$ 
  // Build the count array
2 for  $i \leftarrow 0$  to  $p - 1$  do
3    $C_{m_i} \leftarrow C_{m_i} + 1$ 
4 end
  // Compute the running sum
5 for  $i \leftarrow 1$  to  $p - 1$  do
6    $C_i \leftarrow C_i + C_{i-1}$ 
7 end
8  $R \leftarrow [0, \dots, 0]$ 
  // Reconstruct the sorted array
9 for  $i \leftarrow p - 1$  to  $0$  do
10   $C_{m_i} \leftarrow C_{m_i} - 1$ 
11   $R_{C_{m_i}} \leftarrow m_i$ 
12 end
13 return  $R$ 
```

Figure 8.1: Image taken from the paper

Cleartext implementation of the Counting sort:

```
def counting_sort(a, rng):
    """
    This function implements the counting sort. It takes as input the array a (which is assumed to
    ↪ only have positive values), and a rng (range), which is the 'length-of-the-range' from where
    ↪ the values in a come from.
    """
    # build the count array
    C = np.zeros(shape=rng, dtype=int)
    for i in range(a.shape[0]):
        mi = a[i]
        C[mi] = C[mi] + 1

    # calculating the running sum
    for i in range(1, C.shape[0]):
        C[i] = C[i] + C[i-1]

    # reconstructing the sorted array
    R = np.zeros(a.shape[0], dtype=int)
    for i in range(a.shape[0]-1, -1, -1):
        mi = a[i]
        C[mi] = C[mi] - 1
        R[C[mi]] = mi

    return R
```

8.1.1 Blind Counting Sort

The blind counting sort is counting sort implemented for ciphertexts. The way it works is that works on a RLWE ciphertext that contains the numbers you want to sort as coefficients, which are then sorted by changing the positions of the coefficients in the RLWE ciphertext, the paper calls this RLWE ciphertext a LUT (Lookup Table). However, this is something im a little iffy about doing because first, Ive seen LUT mostly being used when talking about PBS, and second, LUT in PBS is trivially encrypted, while here, it may or may not be the case (most likely it isnt). So to avoid this confusion/inconsistency in the document, Im going to use RLWE and not LUT. The paper implements the following 'access' methods on which the blind sorting algorithm is built.

Blind Array Access (BAA)

Blind Array Access allows one to access a coefficient at a specific position in the RLWE ciphertext. This is done the same way it was being done in PBS, you first rotate the ciphertext such that the coefficient that you want is in the front, and then use the SampleExtract method to extract the 0th coefficient. This is the Blind reading operation.

Algorithm 2: Blind Array Access (BAA)	
Input	An encrypted index $\llbracket i \rrbracket_{\text{LWE}}$ A LUT ciphertext $\llbracket m_0, \dots, m_{p-1} \rrbracket_{\text{LUT}}$
Output	A LWE ciphertext $\llbracket m_i \rrbracket_{\text{LWE}}$
1	$\llbracket \text{rotated} \rrbracket_{\text{LUT}} \leftarrow BR(\llbracket i \rrbracket_{\text{LWE}}, \llbracket m_0, \dots, m_{p-1} \rrbracket_{\text{LUT}})$
2	$\llbracket m_i \rrbracket_{\text{LWE}} \leftarrow SE(0, \llbracket \text{rotated} \rrbracket_{\text{LUT}})$
3	return $\llbracket m_i \rrbracket_{\text{LWE}}$

Figure 8.2: Image taken from the paper

Blind Array Add (BAAdd)

Blind Array Add is the blind writing operation. Given the RLWE ciphertext, this operation adds an LWE ciphertext to a specific coefficient of the RLWE ciphertext that is also provided.

Algorithm 3: Blind Array Add (BAAdd)

Input : An encrypted index $\llbracket i \rrbracket_{\text{LWE}}$
A LUT ciphertext $\llbracket m \rrbracket_{\text{LUT}}$
An encrypted value $\llbracket x \rrbracket_{\text{LWE}}$
Output: A LUT ciphertext $\llbracket m + x\delta_i \rrbracket_{\text{LUT}}$

- 1 $\llbracket x\delta_0 \rrbracket_{\text{LUT}} \leftarrow \text{PFKS}(\llbracket x \rrbracket_{\text{LWE}})$
- 2 $\llbracket x\delta_i \rrbracket_{\text{LUT}} \leftarrow \text{BR}(-\llbracket i \rrbracket_{\text{LWE}}, \llbracket x\delta_0 \rrbracket_{\text{LUT}})$
- 3 **return** $\llbracket m \rrbracket_{\text{LUT}} + \llbracket x\delta_i \rrbracket_{\text{LUT}}$

Figure 8.3: Image taken from the paper

The blind counting sort then uses the operations defined above to then sort the coefficients of the RLWE ciphertext. It basically just mimics the cleartext algorithm except that it uses the respective blind methods.

Algorithm 4: Blind Counting Sort (BCS)

Input : A LUT ciphertext $\llbracket m_0, \dots, m_{p-1} \rrbracket_{\text{LUT}}$
Output: A sorted LUT

- 1 $\llbracket C \rrbracket_{\text{LUT}} \leftarrow [0, \dots, 0]_{\text{LUT}}$
- 2 **for** $i \leftarrow 0$ **to** $p - 1$ **do**
- 3 $\llbracket m_i \rrbracket_{\text{LWE}} \leftarrow \text{BAA}(\llbracket i \rrbracket_{\text{LWE}}, \llbracket m_0, \dots, m_{p-1} \rrbracket_{\text{LUT}})$
 // $C_{m_i} \leftarrow C_{m_i} + 1$
- 4 $\llbracket C \rrbracket_{\text{LUT}} \leftarrow \text{BAAdd}(\llbracket m_i \rrbracket_{\text{LWE}}, \llbracket C \rrbracket_{\text{LUT}}, [1]_{\text{LWE}})$
- 5 **end**
- 6 **for** $i \leftarrow 1$ **to** $p - 1$ **do**
- 7 $\llbracket C_{i-1} \rrbracket_{\text{LWE}} \leftarrow \text{BAA}(\llbracket i - 1 \rrbracket_{\text{LWE}}, \llbracket C \rrbracket_{\text{LUT}})$
 // $C_i \leftarrow C_i + C_{i-1}$
- 8 $\llbracket C \rrbracket_{\text{LUT}} \leftarrow \text{BAAdd}(\llbracket i \rrbracket_{\text{LWE}}, \llbracket C \rrbracket_{\text{LUT}}, \llbracket C_{i-1} \rrbracket_{\text{LWE}})$
- 9 **end**
- 10 $\llbracket R \rrbracket_{\text{LUT}} \leftarrow [0, \dots, 0]_{\text{LUT}}$
- 11 **for** $i \leftarrow p - 1$ **to** 0 **do**
- 12 $\llbracket m_i \rrbracket_{\text{LWE}} \leftarrow \text{BAA}(\llbracket i \rrbracket_{\text{LWE}}, \llbracket m_0, \dots, m_{p-1} \rrbracket_{\text{LUT}})$
 // $C_{m_i} \leftarrow C_{m_i} - 1$
- 13 $\llbracket C \rrbracket_{\text{LUT}} \leftarrow \text{BAAdd}(\llbracket m_i \rrbracket_{\text{LWE}}, \llbracket C \rrbracket_{\text{LUT}}, [-1]_{\text{LWE}})$
 // $R_{C_{m_i}} \leftarrow m_i$
- 14 $\llbracket C_{m_i} \rrbracket_{\text{LWE}} \leftarrow \text{BAA}(\llbracket m_i \rrbracket_{\text{LWE}}, \llbracket C \rrbracket_{\text{LUT}})$
- 15 $\llbracket R \rrbracket_{\text{LUT}} \leftarrow \text{BAAdd}(\llbracket C_{m_i} \rrbracket_{\text{LWE}}, \llbracket R \rrbracket_{\text{LUT}}, \llbracket m_i \rrbracket_{\text{LWE}})$
- 16 **end**
- 17 **return** $\llbracket R \rrbracket_{\text{LUT}}$

Figure 8.4: Image taken from the paper.

8.1.2 Blind Top-k

The Blind-top k finds the k smallest values from the given set of LWE Ciphertexts. To do so, it uses the blind counting sort algorithm, and then obtains the first k elements. However, in case the sequence length is longer than the processing limit of blind counting sort (which is dictated by the plaintext modulus), the sequence is split and a tournament based method of the latter method is used.

8.1.3 How to use this paper to find max

The counting sort sorts the coefficients in the increasing order. A BlindRotate coupled with a SampleExtract (essentially the BAA) operation can be used to then get the last coefficient. However, there is a problem which may or may not be a deal-breaker for this approach. Traditionally, counting sort only works with positive integers (integers is okay, but only positive is not), atleast in cleartext version. The blind counting sort also seems to work only with unsigned (only positive) integers, the way its implemented. However, it might still be possible to work with this by offsetting the negative values to positive (essentially encoding them such that they loop around the plaintext modulus value), and then after performing the sorting operation, we might need to rotate them by (maybe) $p/2$ and that should work.

8.1.4 Runtime Comparison

The runtime for calculating max for different approaches was checked. Experiment characteristics are the following:

- The maximum value was chosen among 16 values.
- The code was run on local machine (16GB RAM, R9 5900HS).
- No rayon based parallelization was used in either of the methods (except the ones that TFHE does internally).
- The max operation was executed 100 times with each method and the runtimes for each run were recorded.

The following approaches were compared:

- **$O(n^2)$ max:** Each element was compared with each other to find the maximum element. This approach is referred to as the *TFHE-native* approach as it uses the native tfhe-rs methods available to calculate the maximum.
- **Counting Sort based max:** The paper's blind counting sort is used to find the max from the same array. This approach is referred to as the *paper based* approach.




Figure 8.5: Runtime Comparison Boxplot

Approach	Mean Runtime (sec)
TFHE-rs Native	4.12
Counting Sort Based	1.35

8.2 Current Status

- I am now able to run both the blind sort and blind top-k algorithms. [DONE]
- I need to compare the runtime of the current approach (iteratively finding max) and the runtime of this algorithm for finding max. [DONE]
- Ensure that the 'trick' to make it work with negative numbers works. [In Progress]

Chapter 9

References

- Core-crypto - TFHE-rs
- Daniel Lowengrub's Blog on TFHE implementation from scratch.
- The TFHE-rs Handbook.(Read until the PBS section, it describes a lot of other stuff too.)
- TFHE-rs's blog series: TFHE Deep Dive
- Hitchhiker's Guide to the TFHE Scheme (havent gone through yet).
- A non-comparison oblivious sort and its application to private k-NN
- CODE: A non-comparison oblivious sort and its application to private k-NN